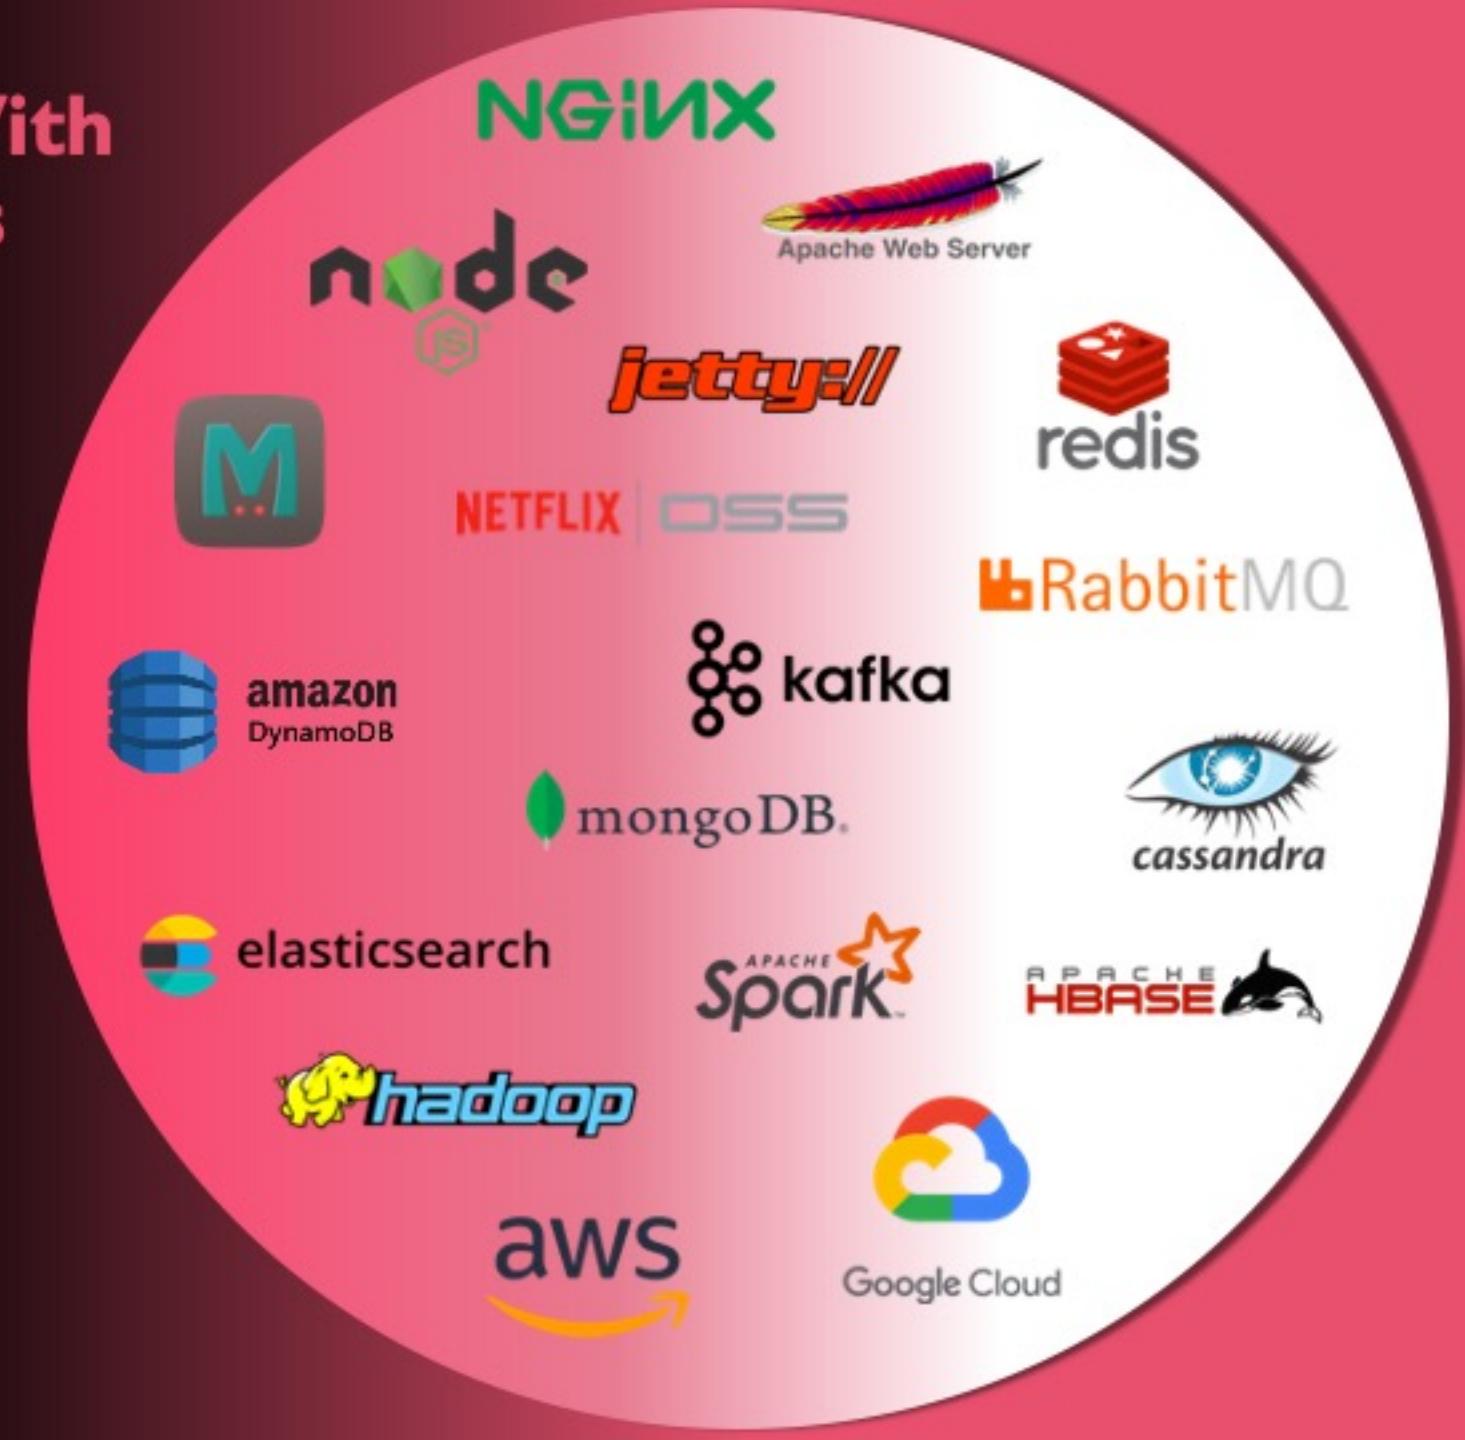
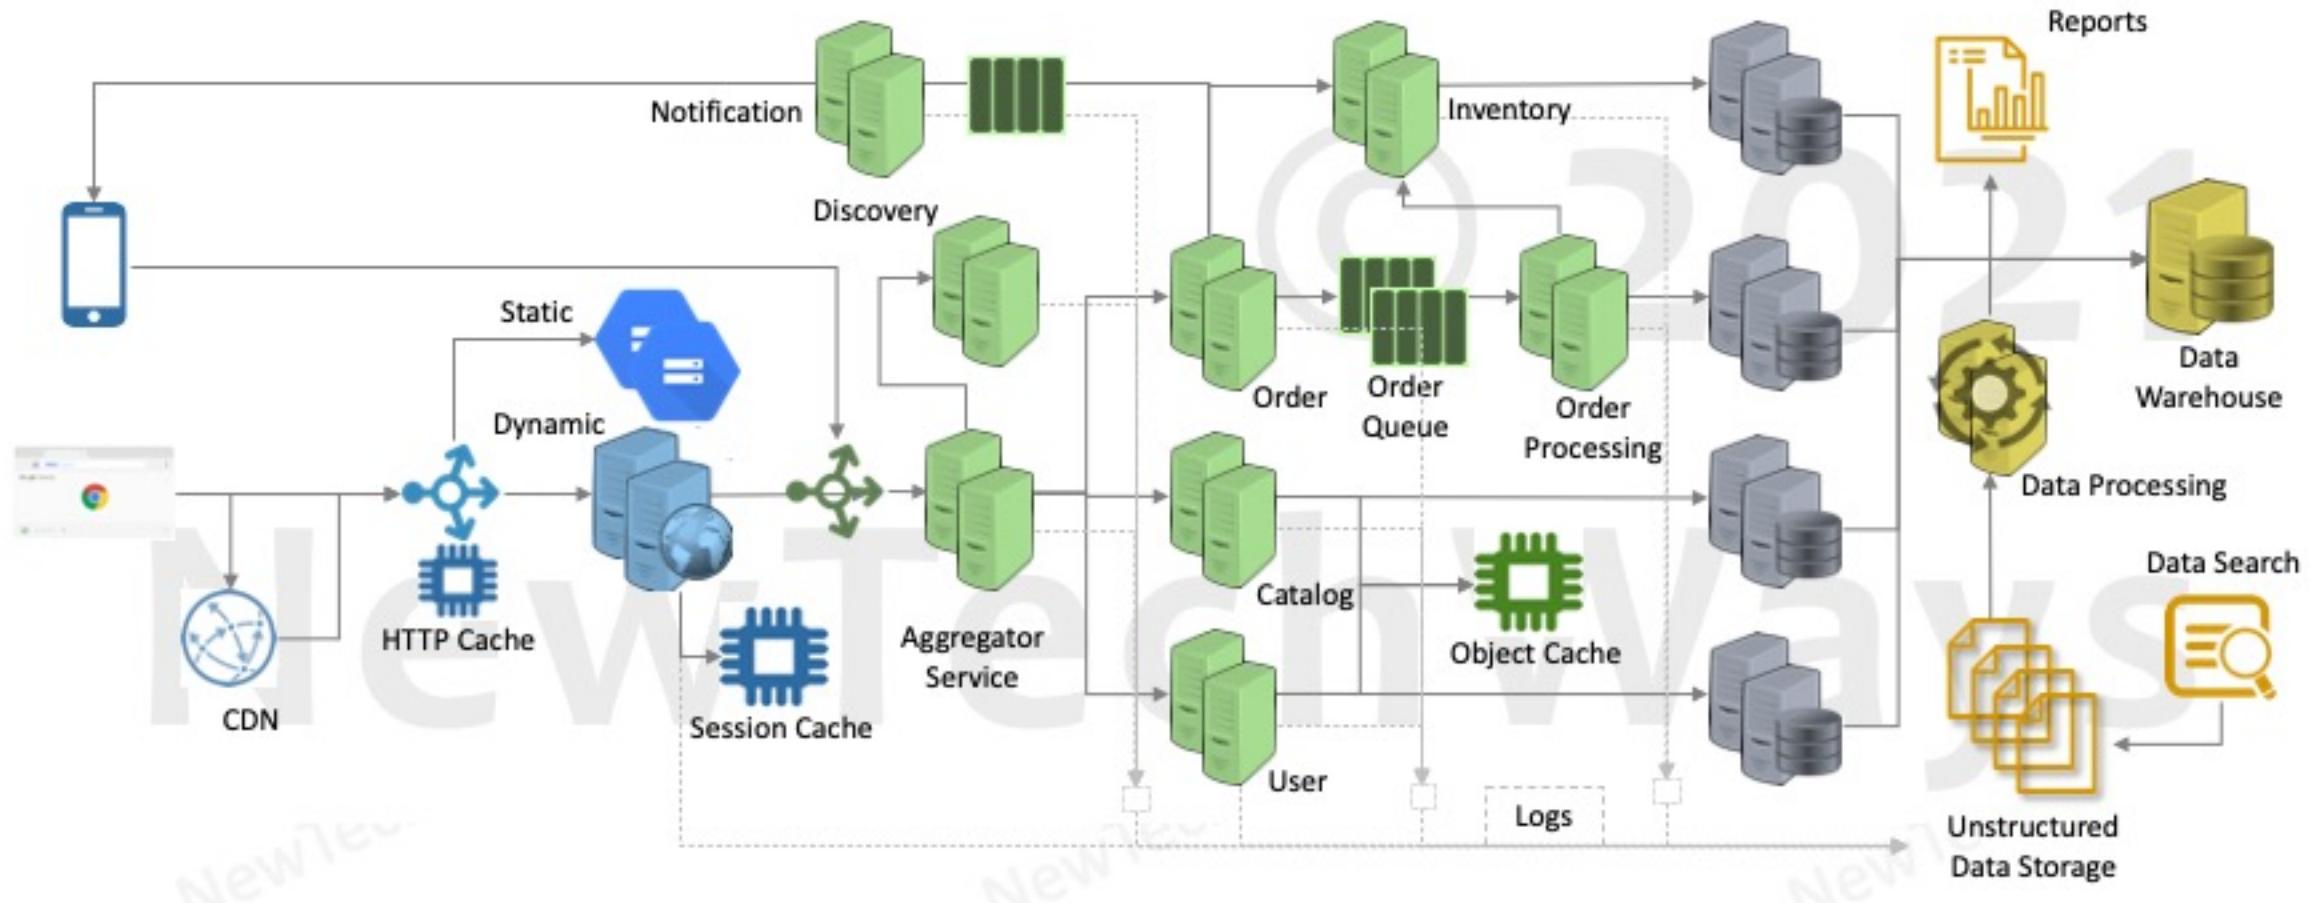


Architecting Solutions With Platforms & Frameworks

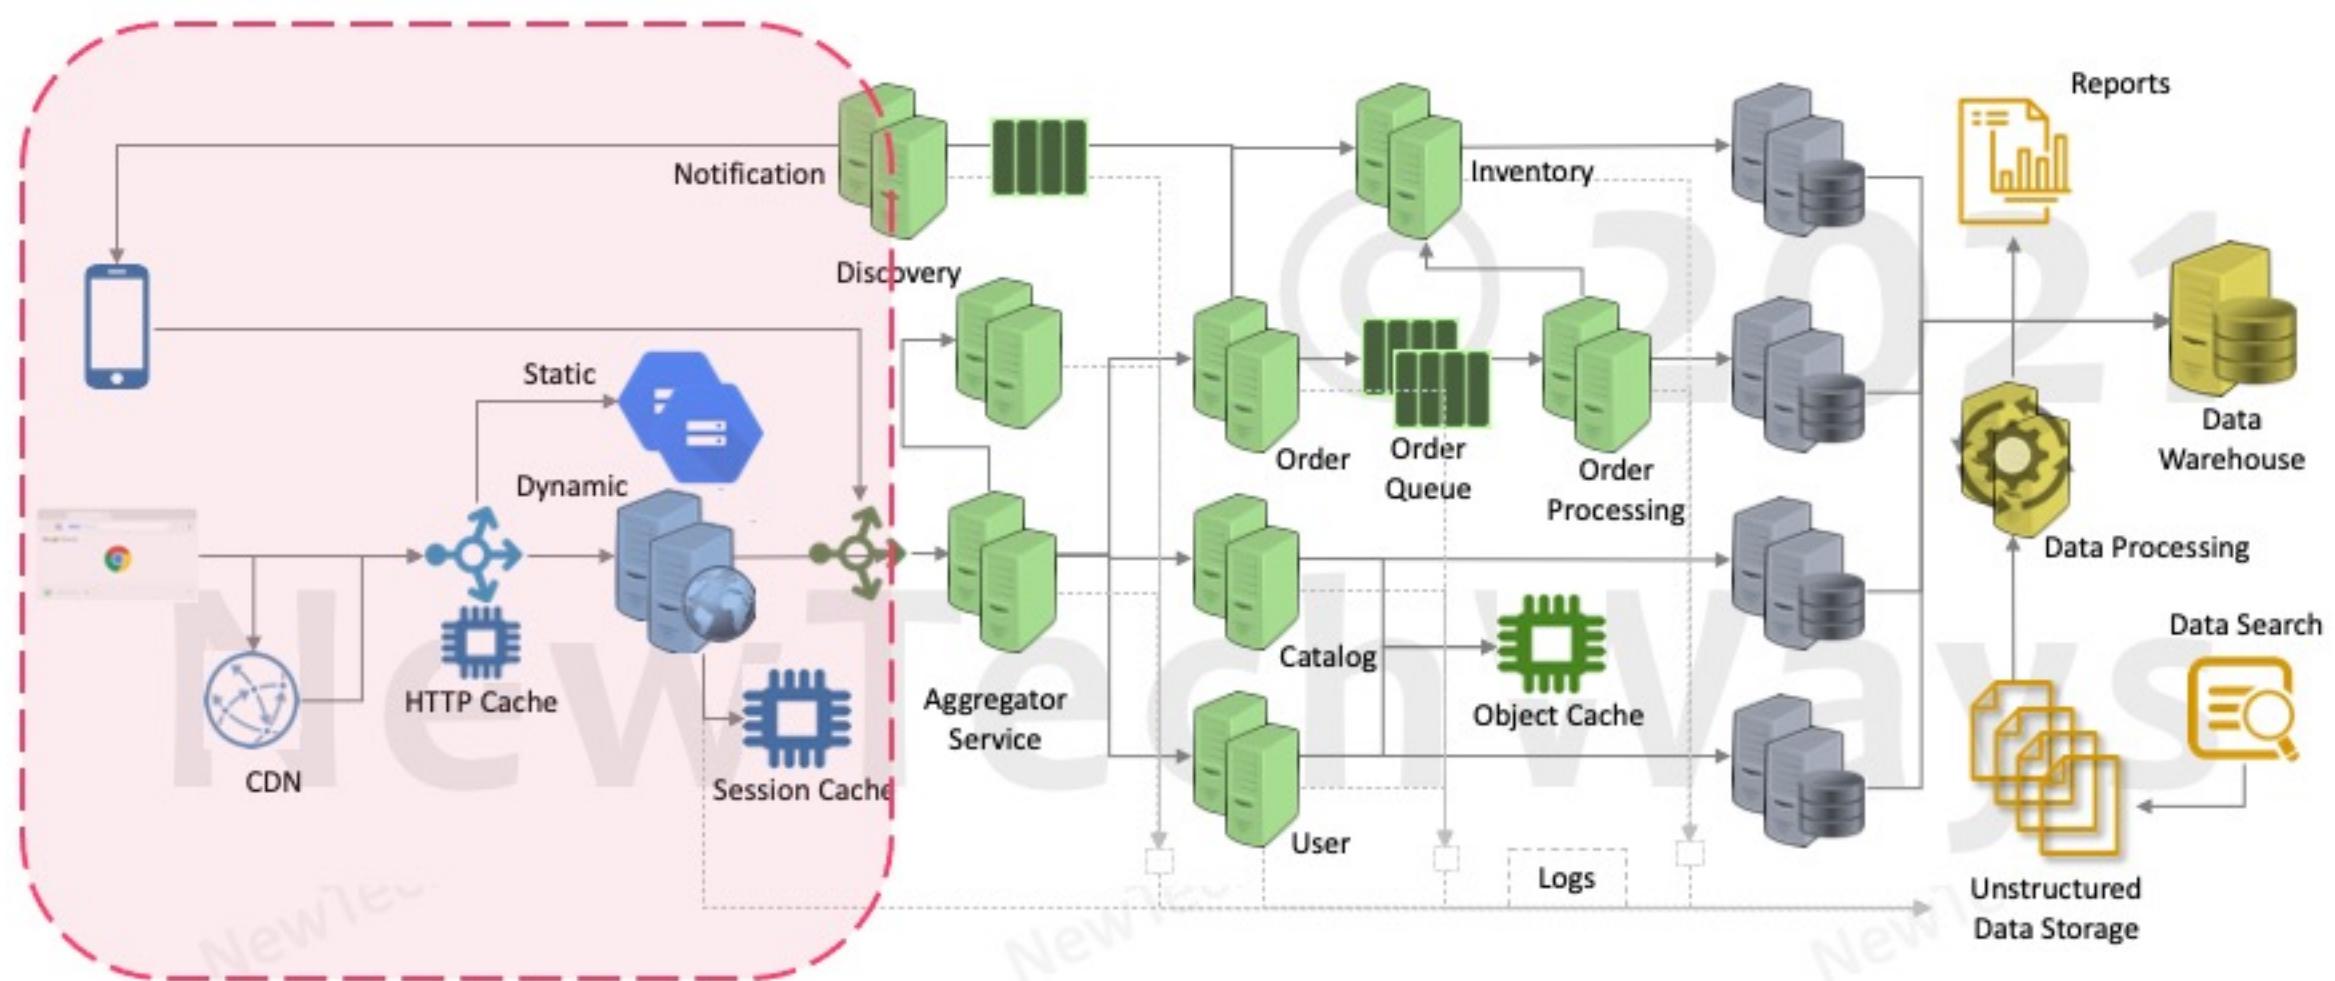
- Platforms for
 - Web Apps
 - Services
 - Datastores
 - Analytics
- Platform Functionality
- Platform Architecture
 - Performance
 - Scalability
 - Reliability
- Platform Use-Cases
- Platform Alternatives
 - Comparison
- Architecting Solution
 - End-to-End



Reference System



Web Applications



Web Application – Solutions

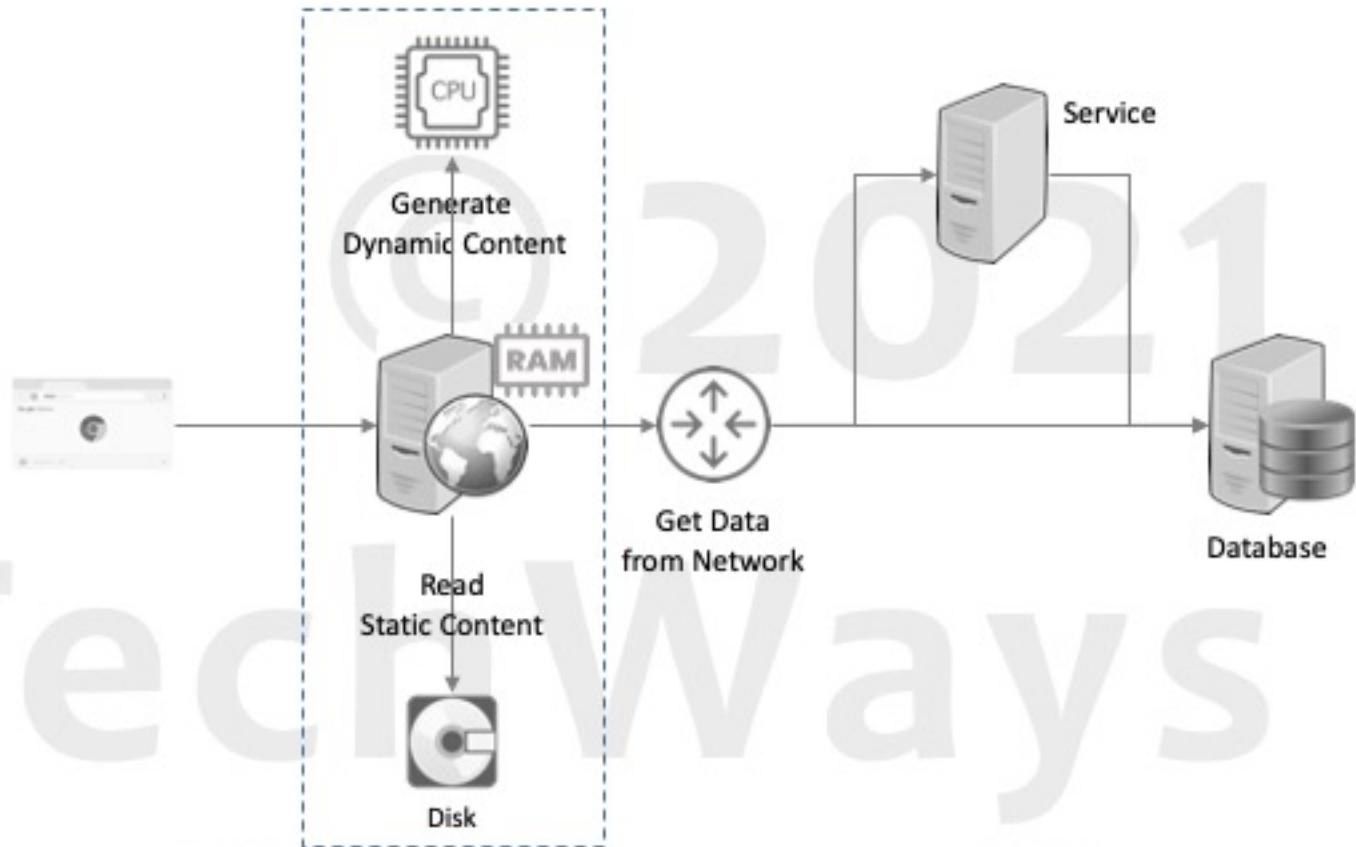
- Static Content
 - Apache Web Server
 - Nginx Web Server
 - Cloud Storage
- Dynamic Content
 - Web Server – Apache HTTPD, NodeJS
 - Java Web Containers – Tomcat, Jetty, Spring-Boot
- Content Caching
 - Nginx
- Content Distribution
 - CDN

© 2021

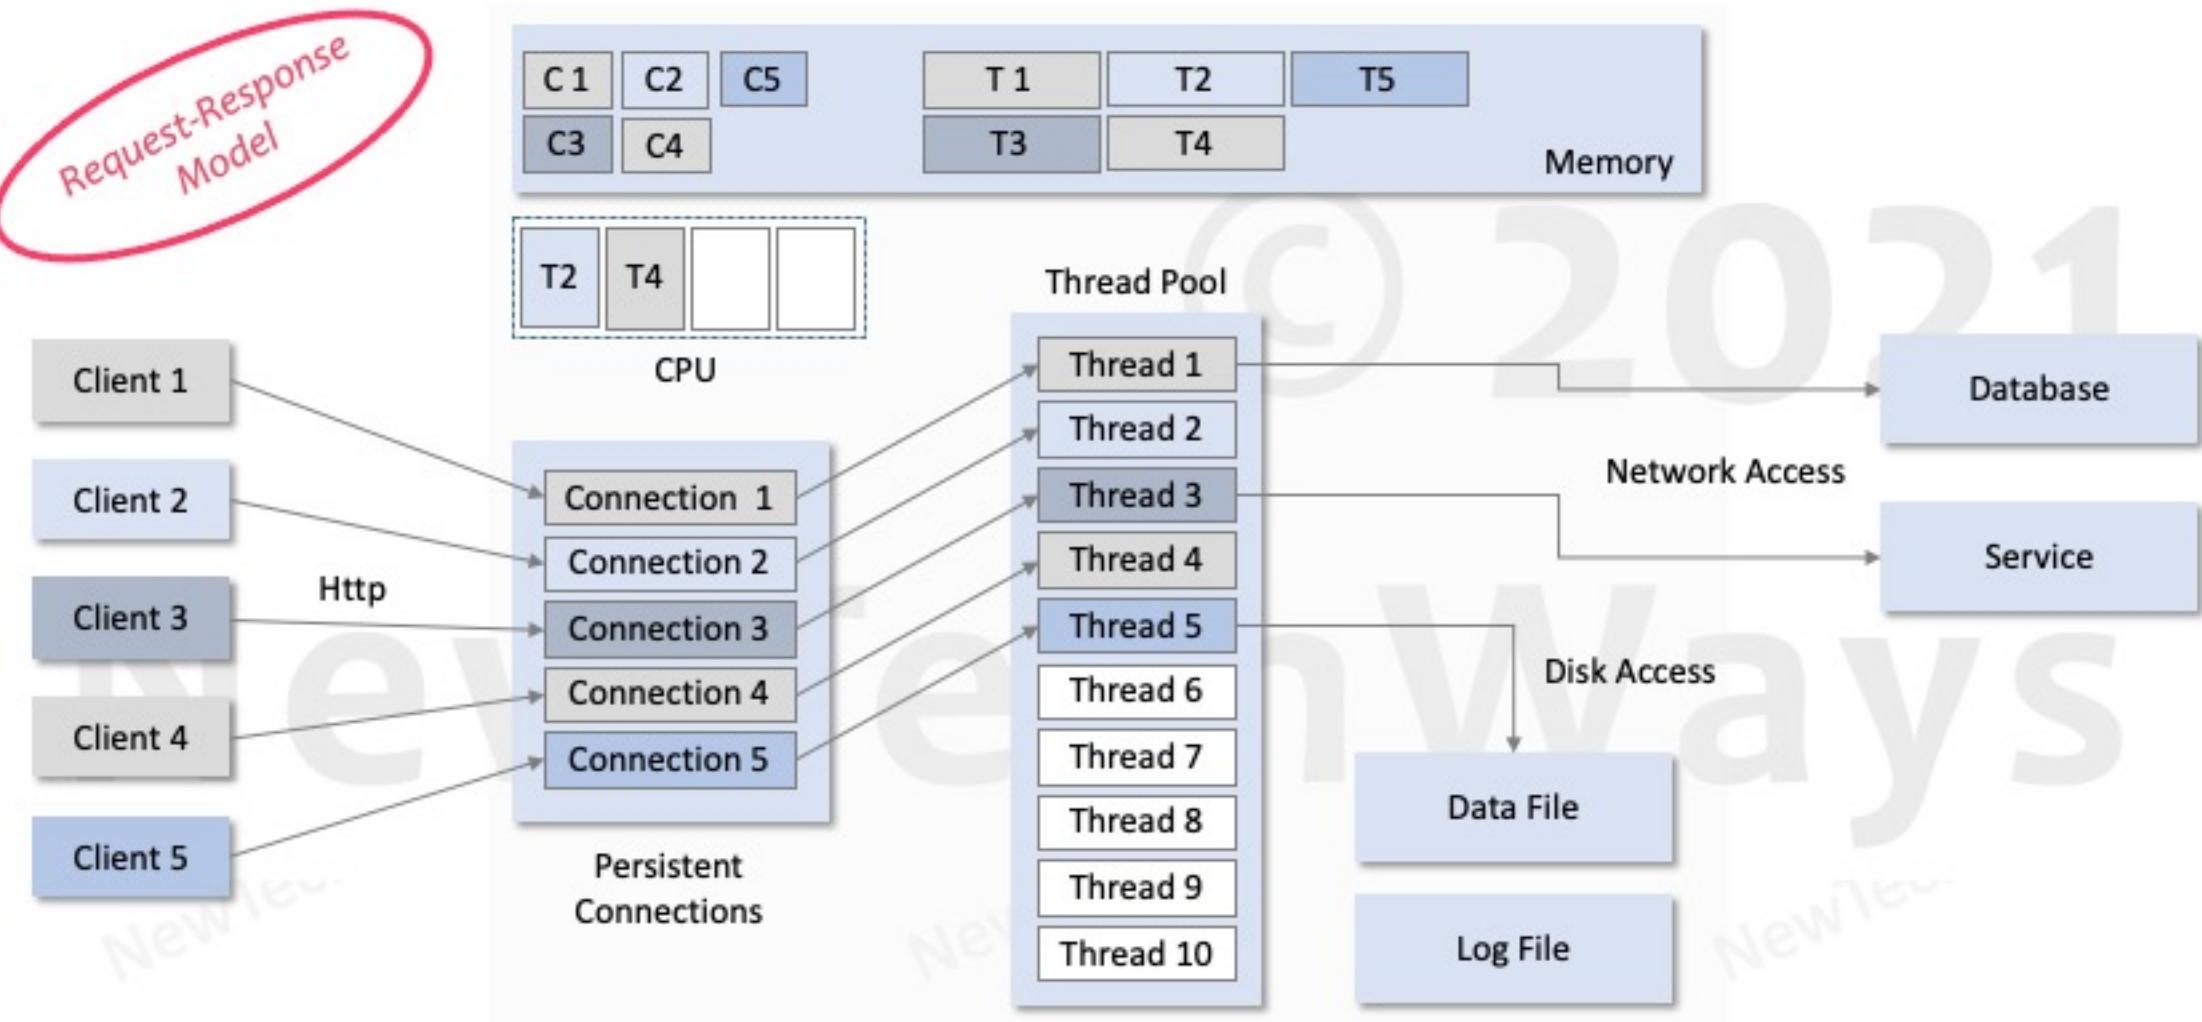
NewTechWays

Apache Webserver

- Store Static Content
 - Html/CSS/JS files
 - Image files
 - Documents
- Generate Dynamic Content
 - Get data & generate pages dynamically
 - PHP, Python, Perl
 - **No JSP/Servlets**
- Act as a Reverse Proxy
 - **Not great**

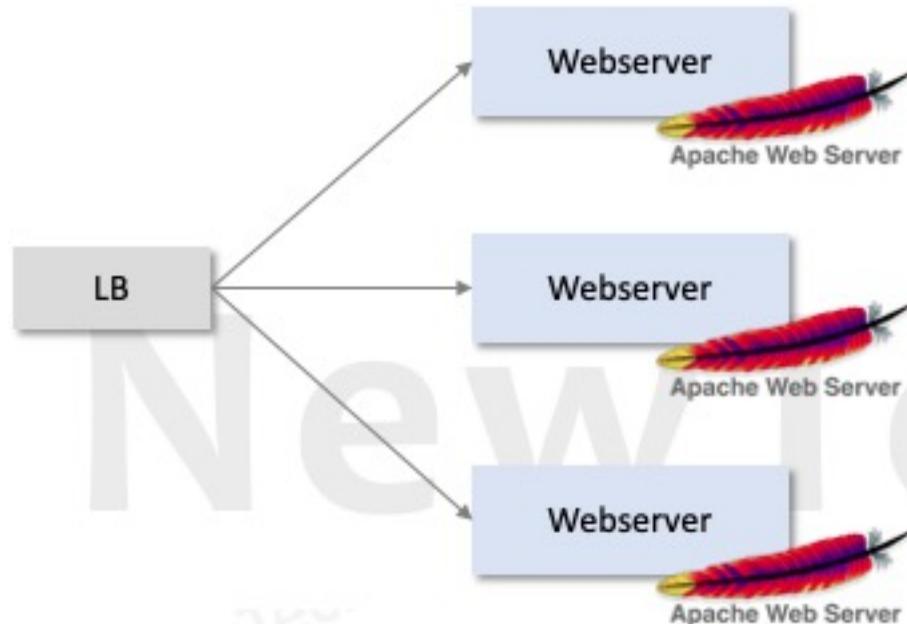


Apache Webserver Architecture

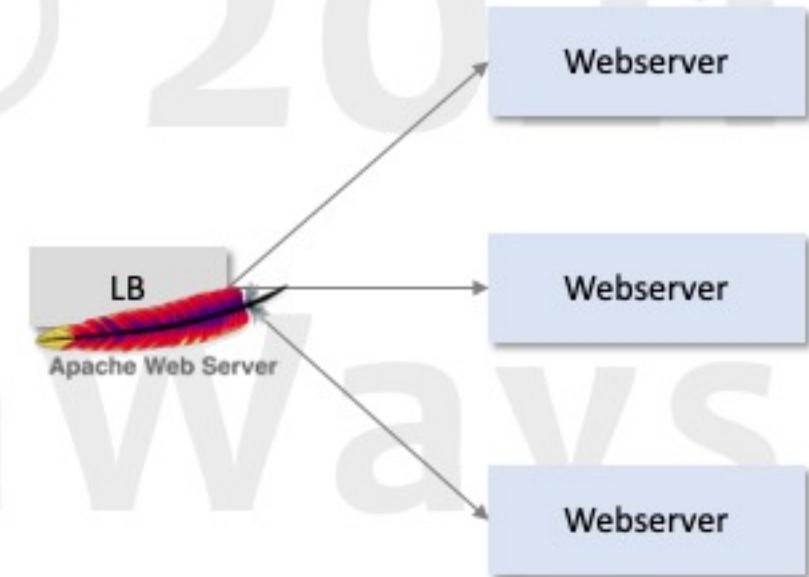


Apache Webserver Scalability

Apache as Webserver
(CPU Bound)

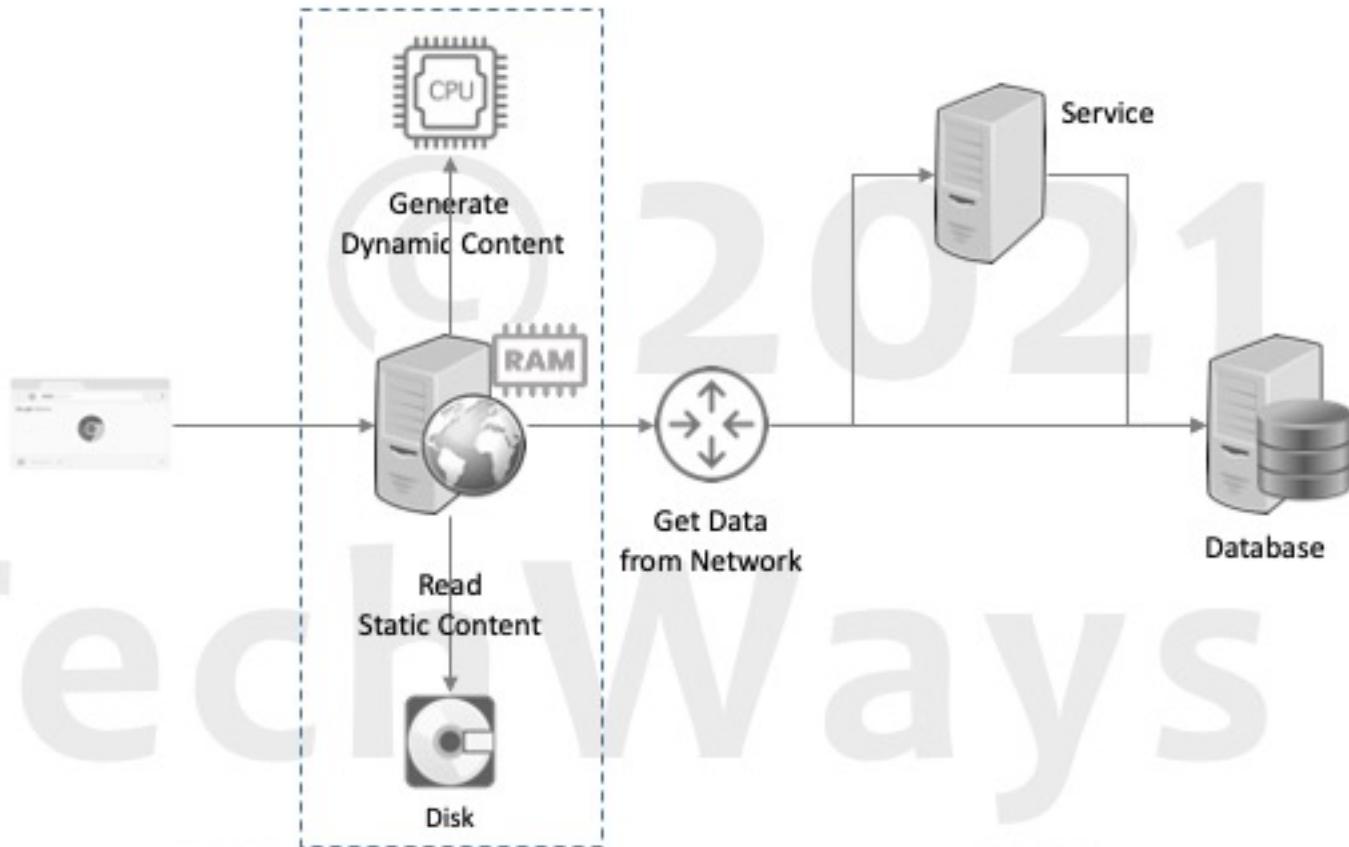


Apache as Reverse Proxy
(IO Bound)

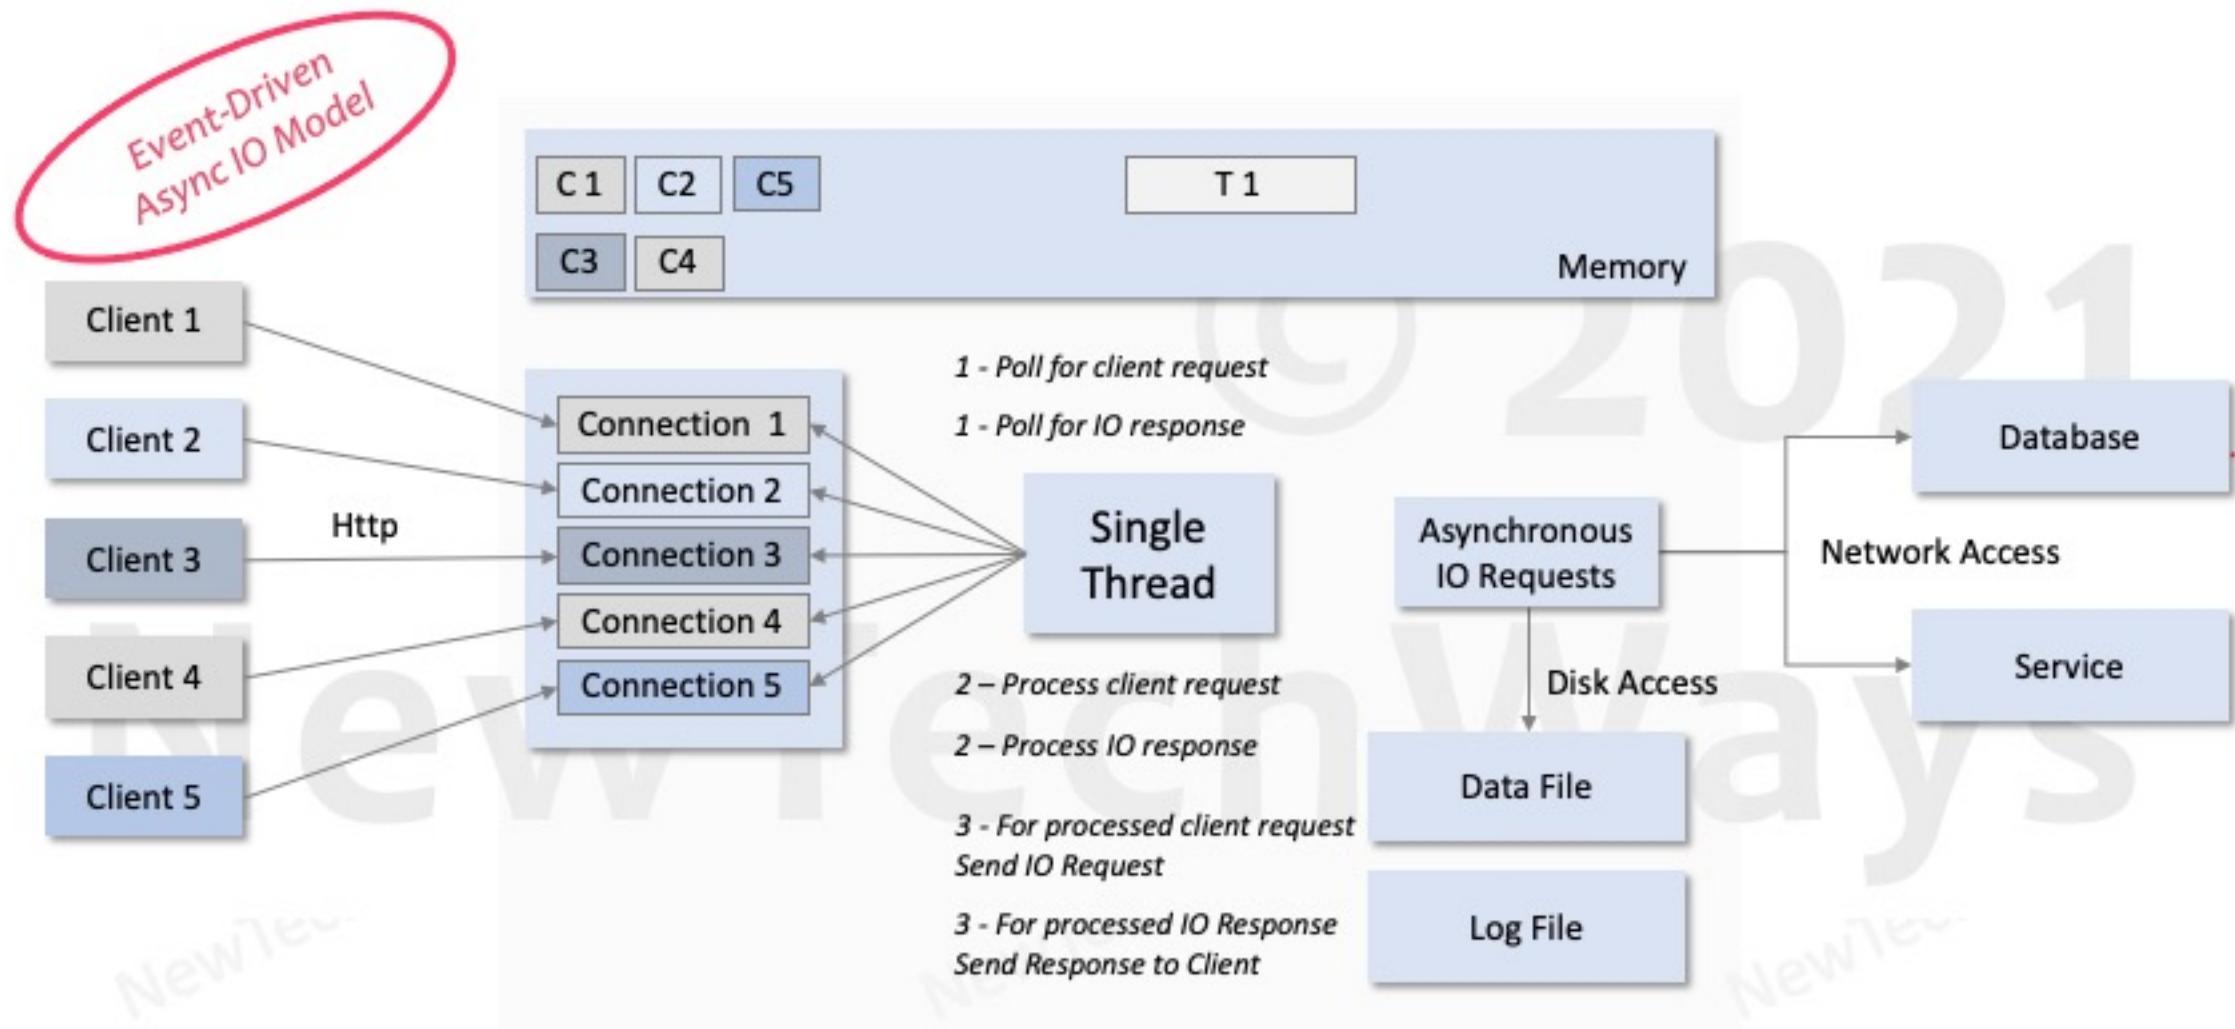


NGINX Webserver

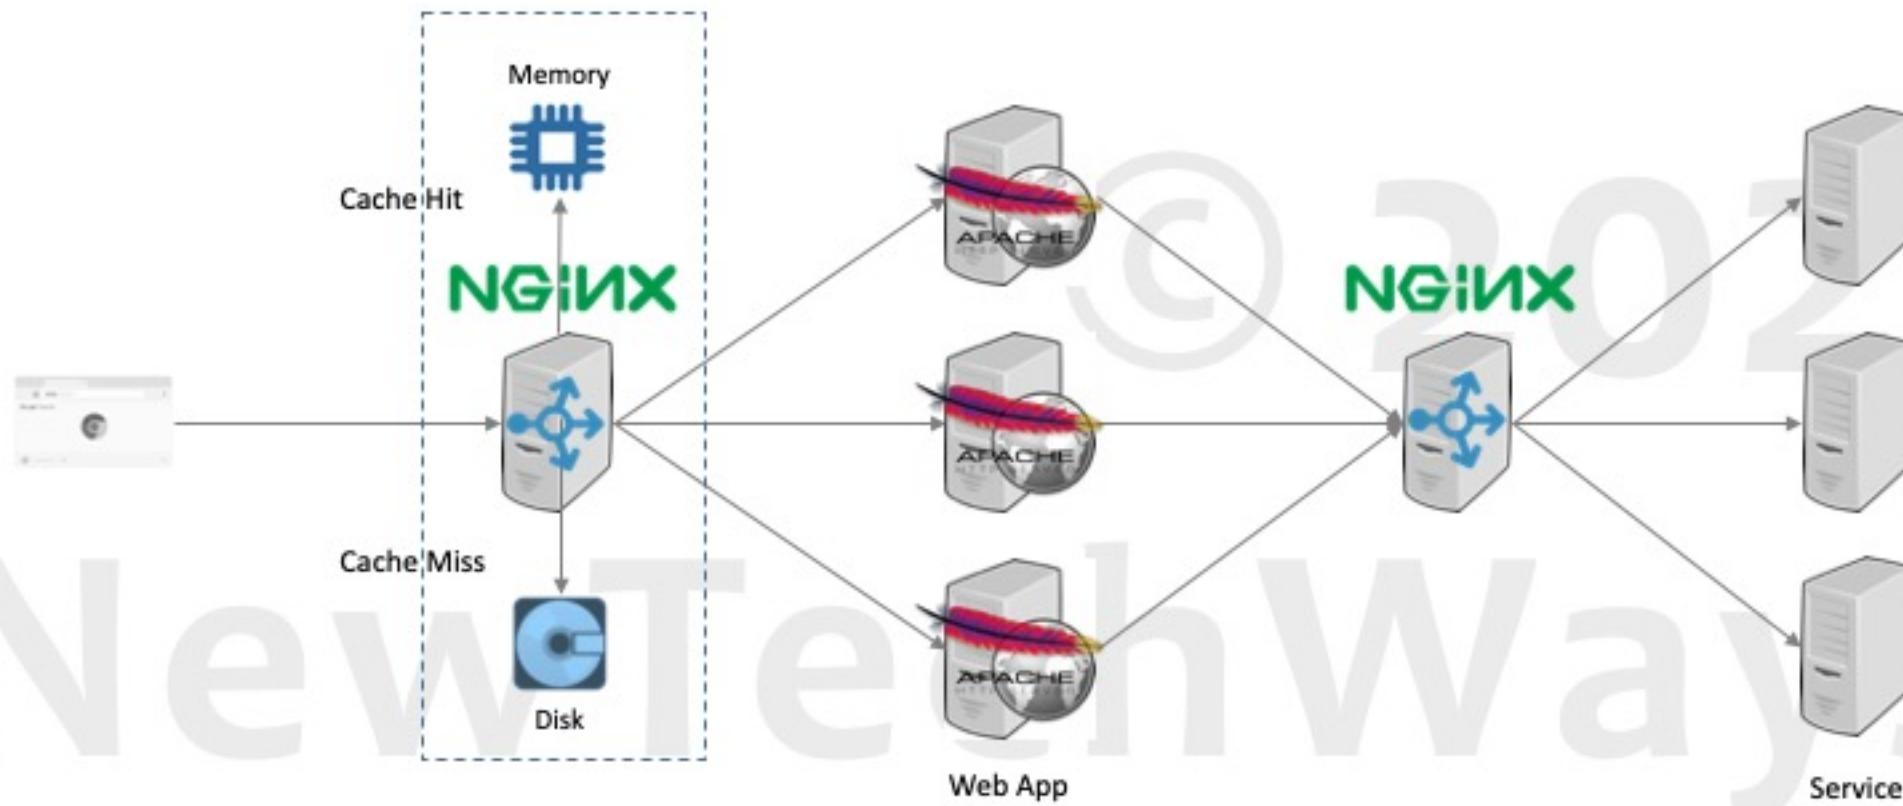
- Store Static Content
 - Html/CSS/JS files
 - Image files
 - Documents
- Generate Dynamic Content
 - Not the best
- Act as a Reverse Proxy
 - Excellent
- Cache Content
 - Good



Nginx Architecture



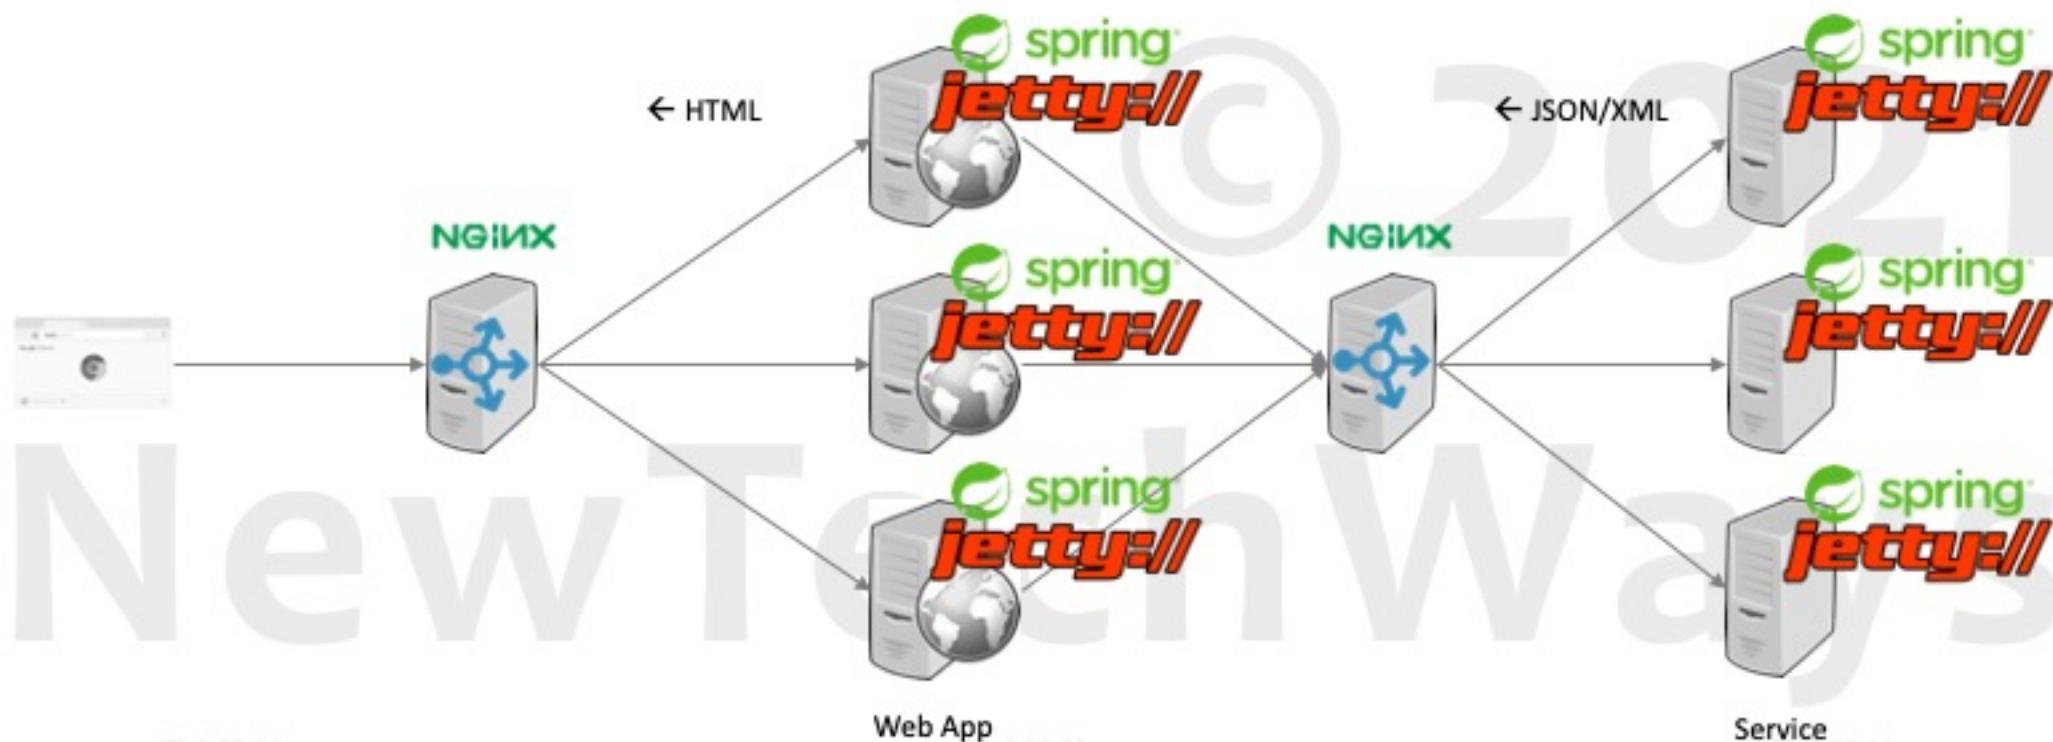
Nginx as Reverse Proxy & Cache



Web Containers & Spring Framework

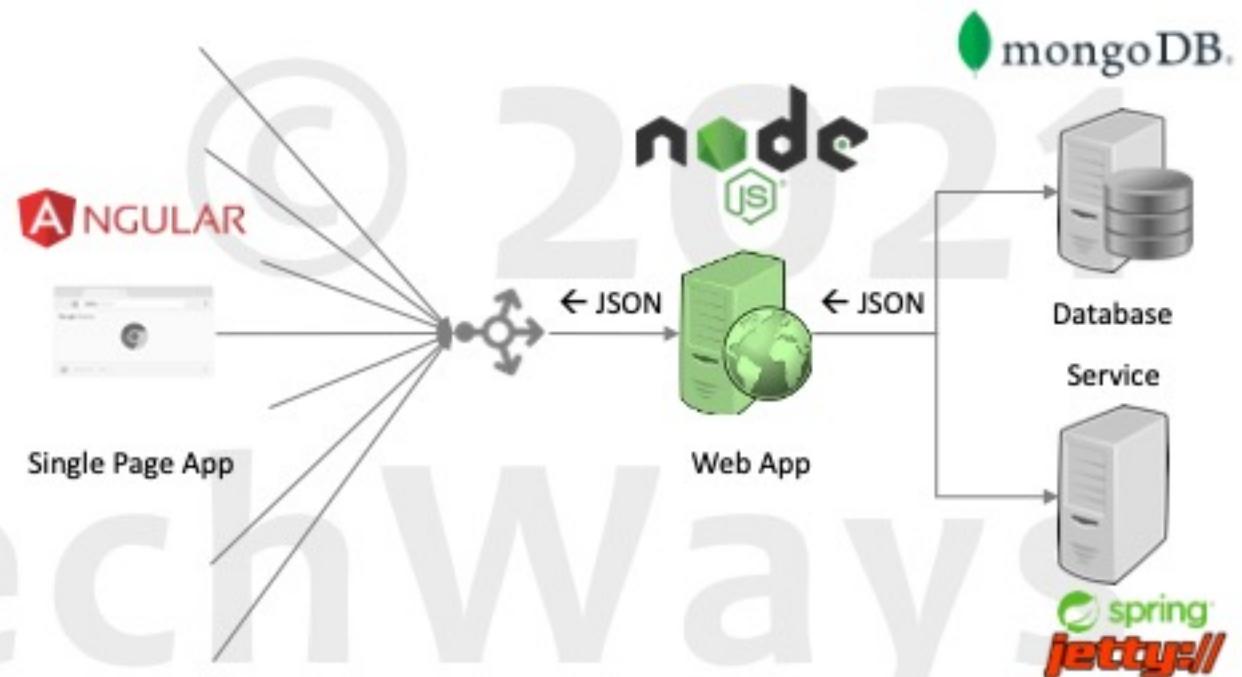
- Dynamic content using Java
 - OO Language for complex logic
 - Web Containers
 - Tomcat
 - Jetty
 - Application Servers
 - Wildfly/JBoss
 - Weblogic
 - Websphere
 - Spring Boot
 - Runs embedded web container
 - Tomcat, Jetty
 - MVC Architecture
 - Servlets for logic
 - JSP for presentation
 - Spring Containers
 - Runs inside a Web container
 - Provides
 - IOC/DI
 - For business logic
 - Model View Controller
 - For frontends
 - JDBC Templates
 - For accessing DB
 - Connection Pools
 - Http, DB
-
- ```
graph LR; subgraph Central [Servlet Engine]; direction TB; SE[Servlet Engine, EJB Container, Session Clustering, Connection Pools, Caching, JMX, JMS, OSGI, ...]; end; AS[Wildfly/JBoss, Weblogic, Websphere] --> Central; SC[IOC/DI, Model View Controller, JDBC Templates, Connection Pools] --> Central;
```

# Jetty & Spring

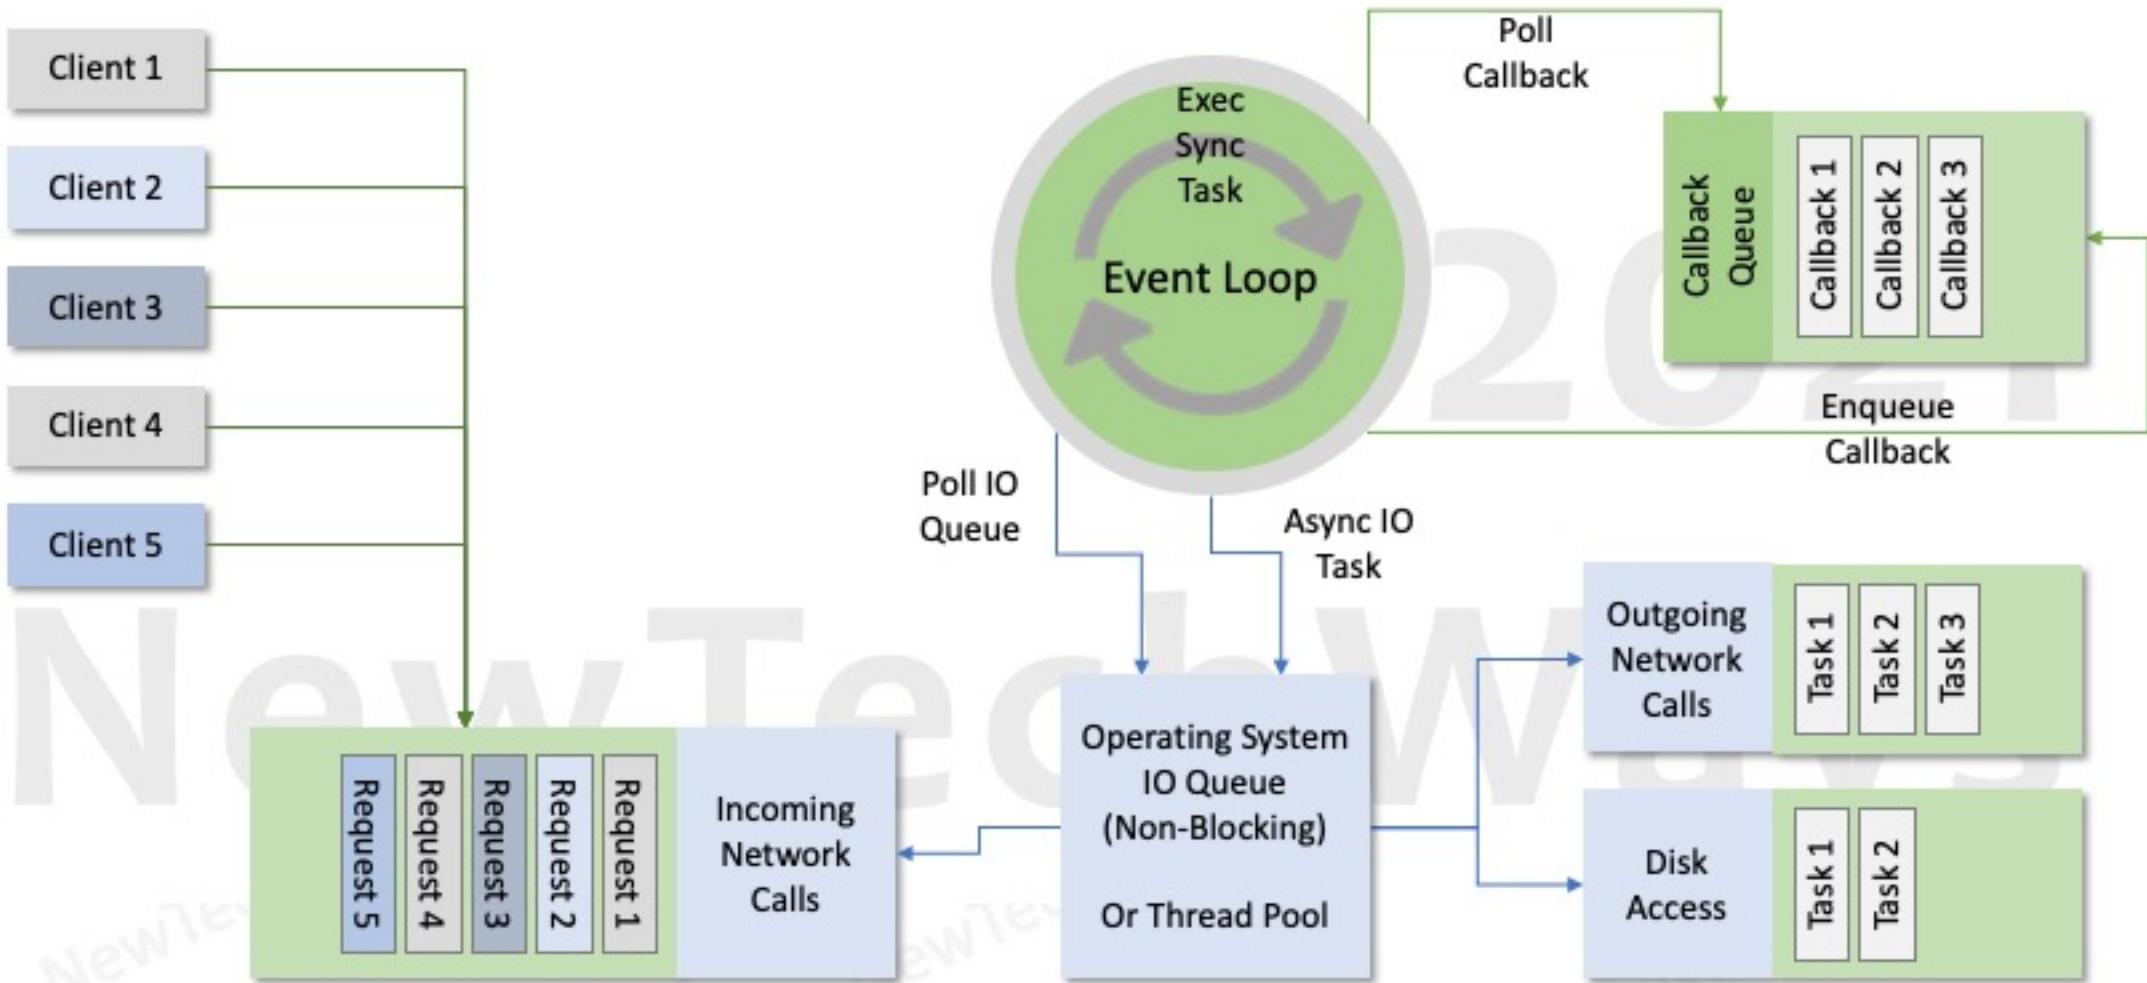


# Node.js

- HTTP server
  - Uses JavaScript engine to process requests
- Highly efficient at handling a large-number of connections
  - For requests that are IO bound and not CPU bound
- Single thread to handle all connections
  - Saves memory
  - Avoids context switching

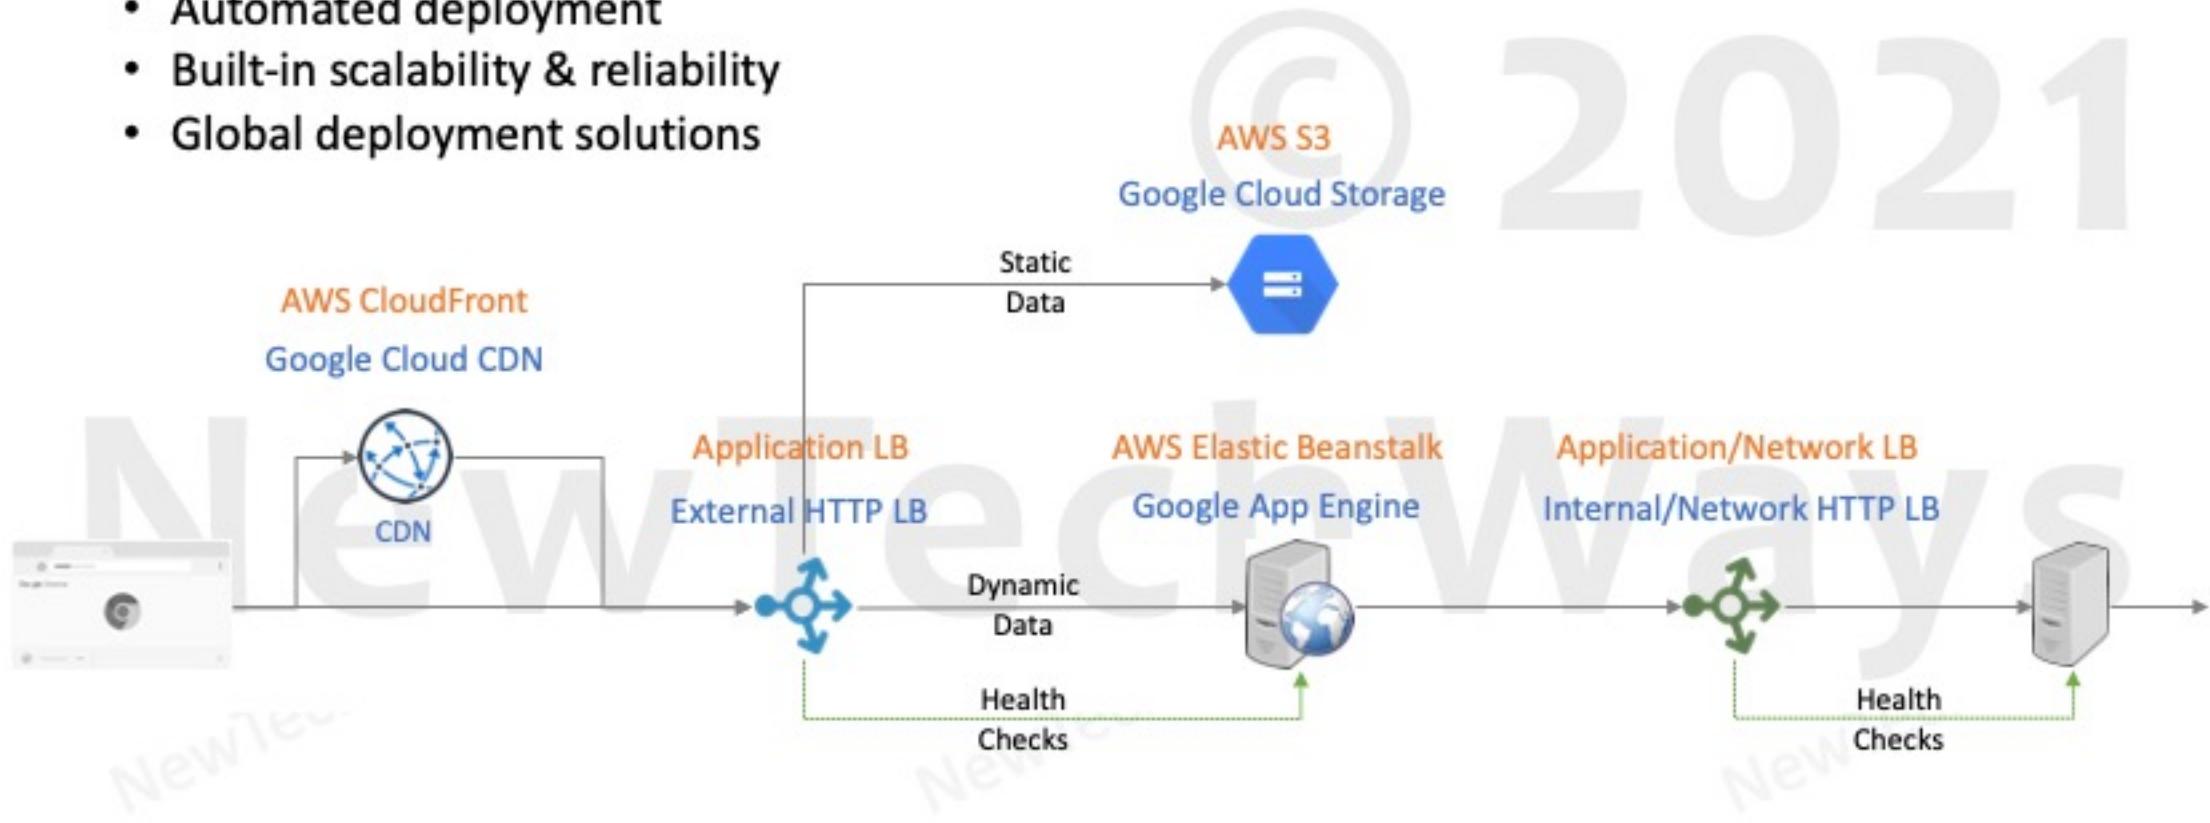


# Node.js – Event Loop



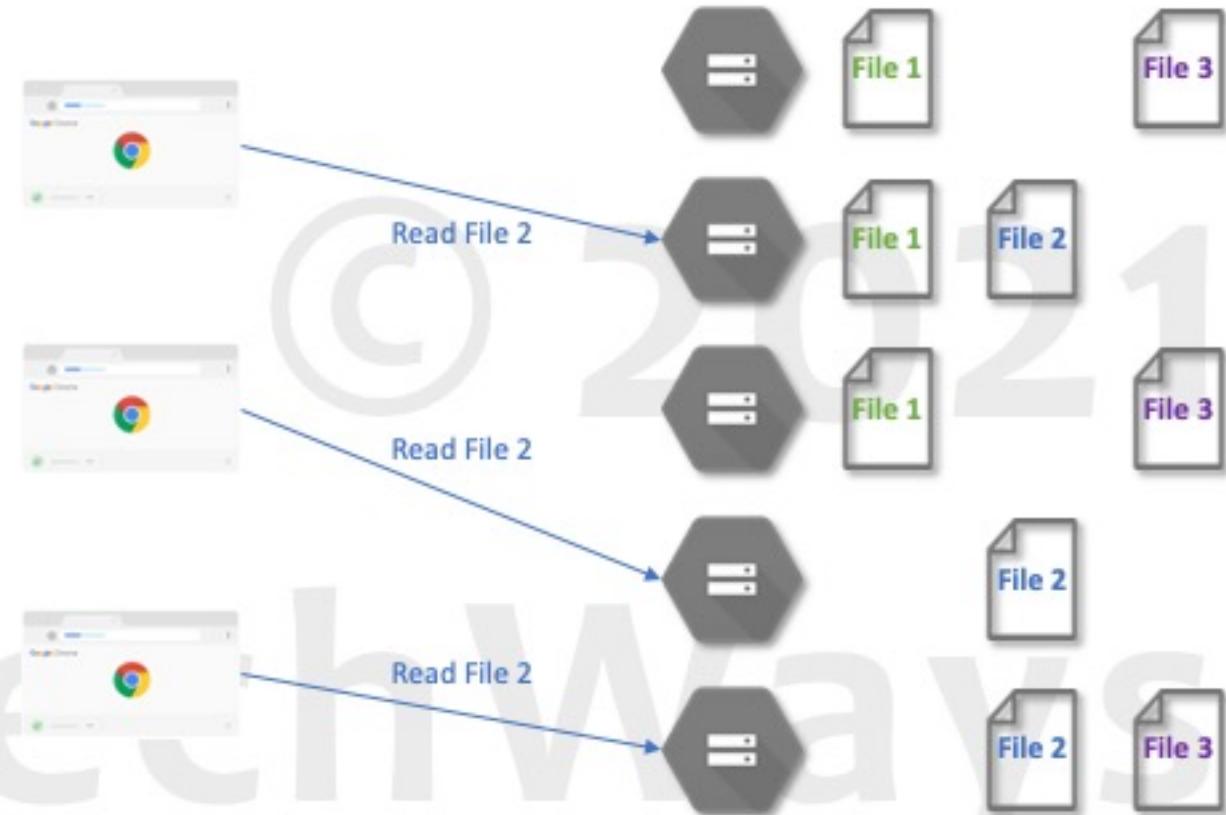
# Cloud Solutions For Web

- Managed Services
  - Hardened solutions
  - Automated deployment
  - Built-in scalability & reliability
  - Global deployment solutions



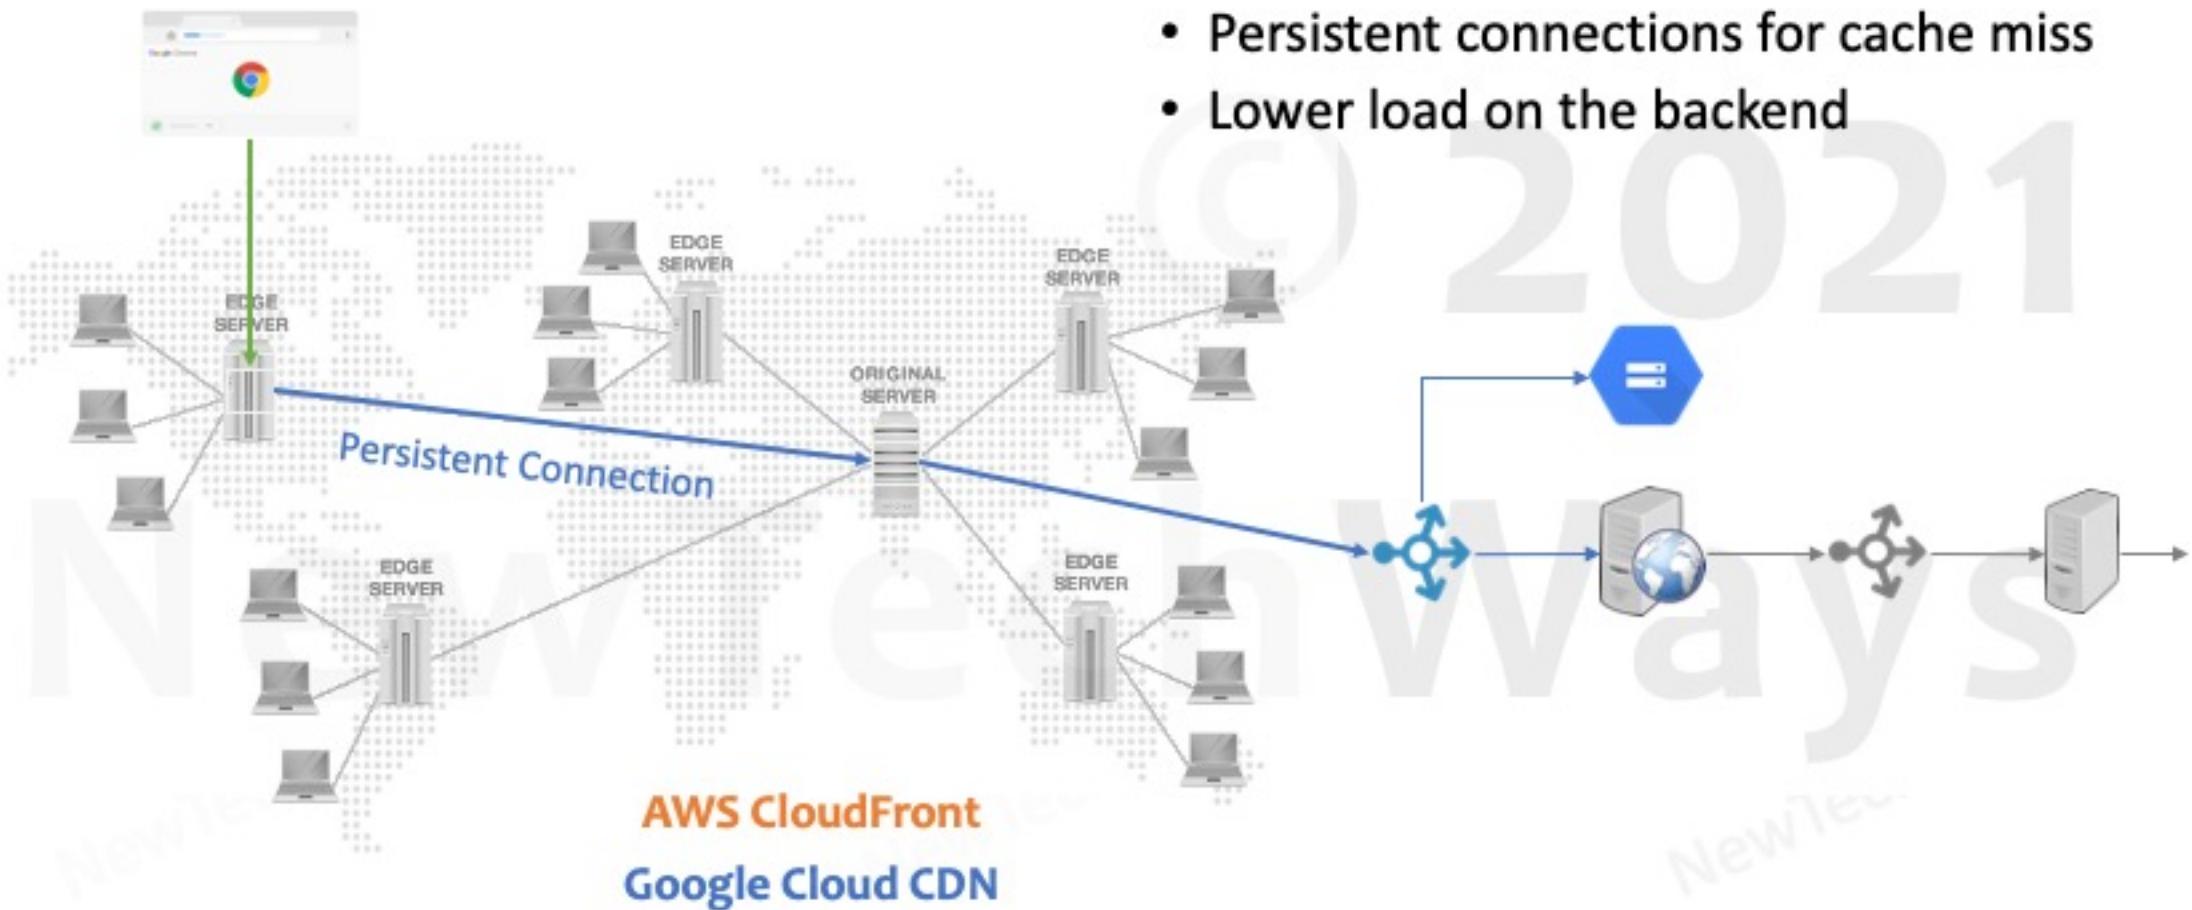
# Cloud Storage

- Unlimited Disk Space
- Version Control
- Access Control
- Low Latency
  - No overhead of Directory structure
- High Throughput
  - Parallel clients
  - Large files can be broken into smaller chunks for parallel read
- High Availability & Reliability
  - Multiple copies
  - Multiple physical locations
- Static Website Creation

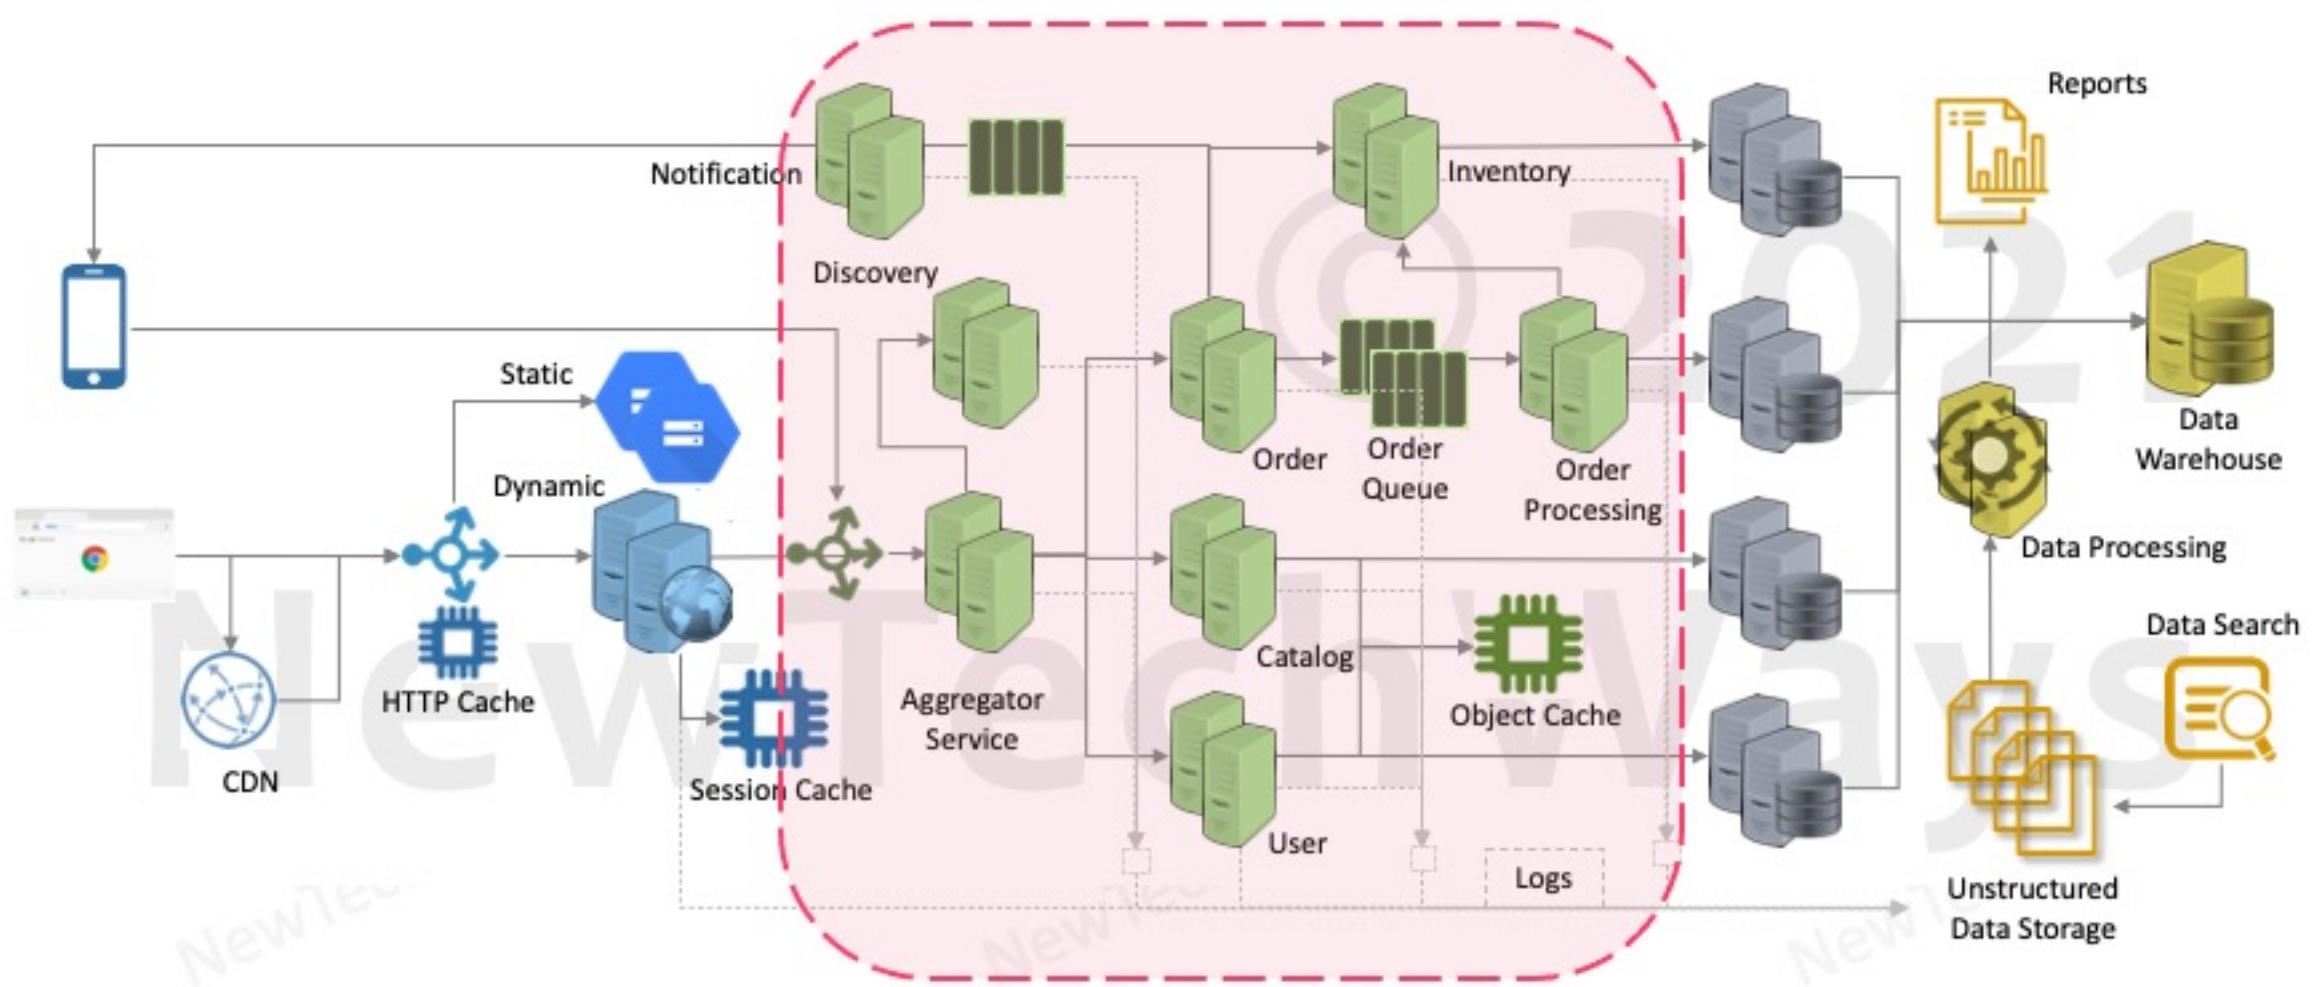


# Cloud CDN

- Low latency local access for cache hits
- Persistent connections for cache miss
- Lower load on the backend



# Services



# Services – Solutions

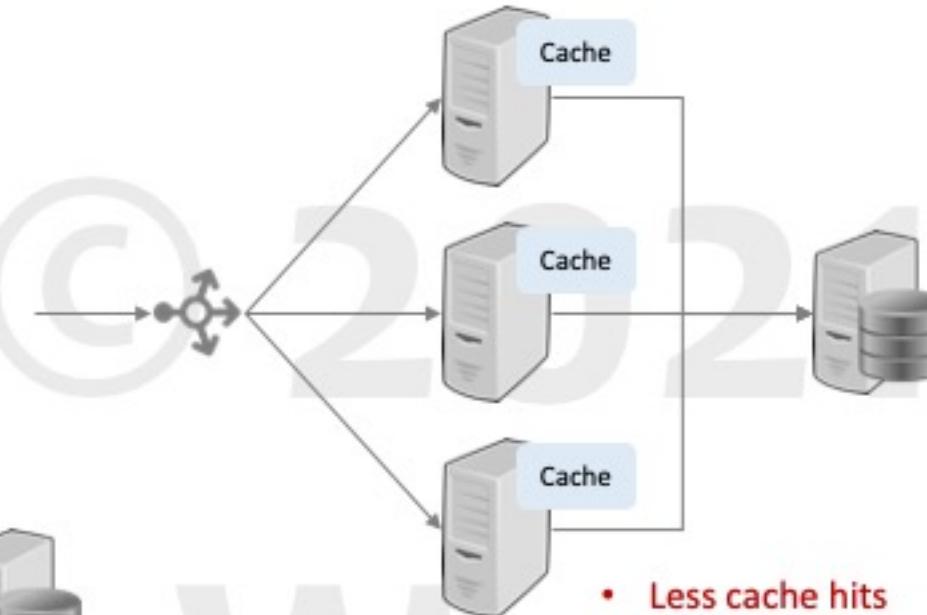
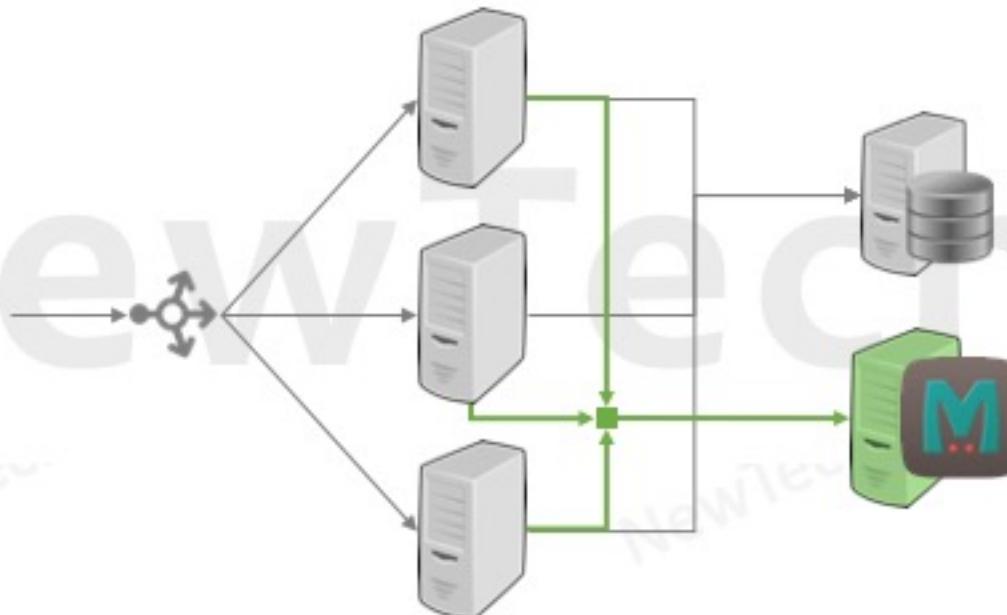
- Web Containers
  - REST & Spring Containers
- Object Caching
  - Memcache
  - Redis
- Asynchronous Messaging
  - Redis
  - RabbitMQ
  - Kafka
- Service Mesh
  - Netflix
  - Istio

© 2021

NewTechWays

# Memcached

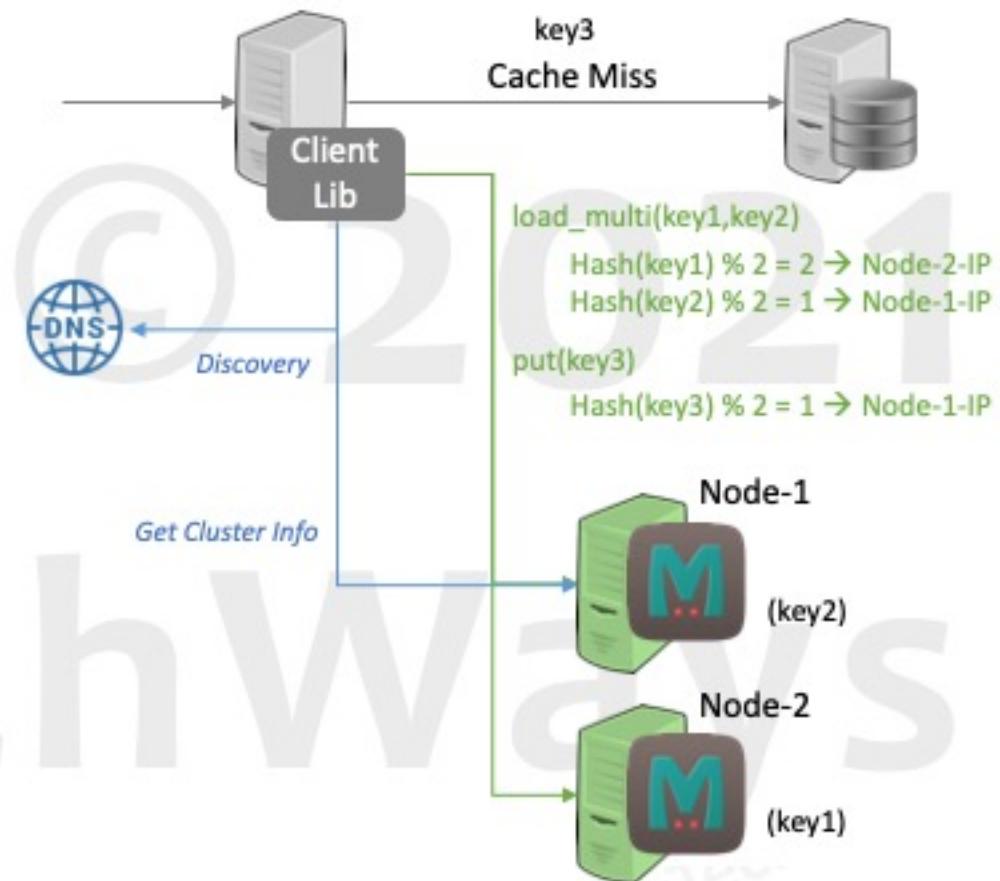
- Stores key value pairs
- Values can be
  - Any blob
  - Any size
    - Preferred < 1 MB
    - Max is configurable



- Less cache hits
- Key-Data duplication
- TTL can be set differently for each data
- Eviction expunges expired data followed by LRU data

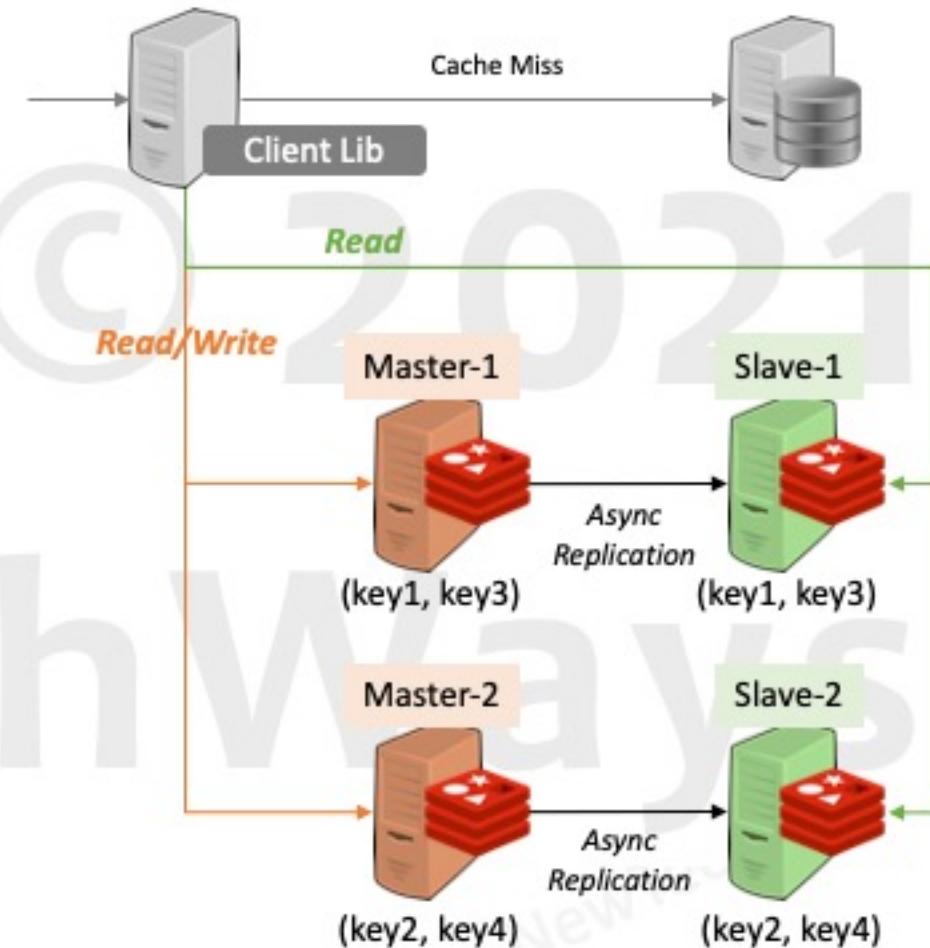
# Memcached Architecture

- Cache-aside pattern
- Sub-millisecond Latency
- Horizontally Scalable
  - Data is partitioned
- High throughput
  - Parallel operations
- Cluster aware client library
  - Consistent hashing for resolving a node
- Node failure is treated as a cache miss
  - Use large number of nodes with less data
- Data is lost if a node crashes or restarted



# Redis Cache

- Much like Memcached
- It is a [ Key → Data structure ] store
  - Strings, Lists, SortedSets, Maps, ...
- Data-Store mode for Persistence
  - Stores data on a disk
  - Allows backups
  - Can be started pre-populated with a backup
- Data Replication
  - Asynchronous and synchronous
  - Read load distribution
  - High Availability
- Can also be used as a messaging queue
- Data-Store mode requires fixed number of nodes
  - Cache mode is suitable for node scaling



# Cloud Caching Solutions

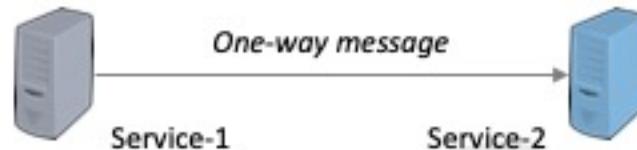
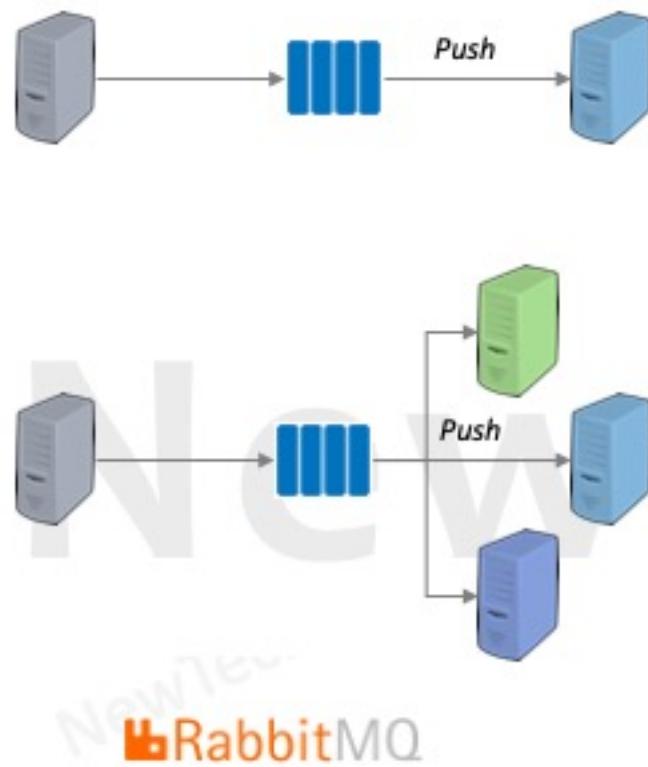
- AWS Elastic Cache
  - Memcached
  - Redis
- Google Memorystore
  - Memcached
  - Redis



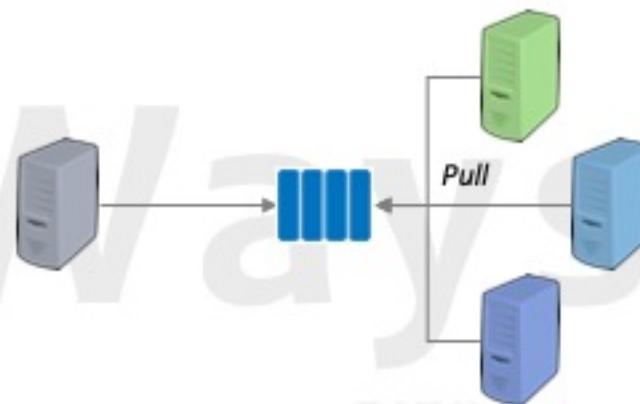
2021

- Fully Managed
- Sub-millisecond latency
- Scalable to 5 TB
- Highly available (99.9%)

# RabbitMQ

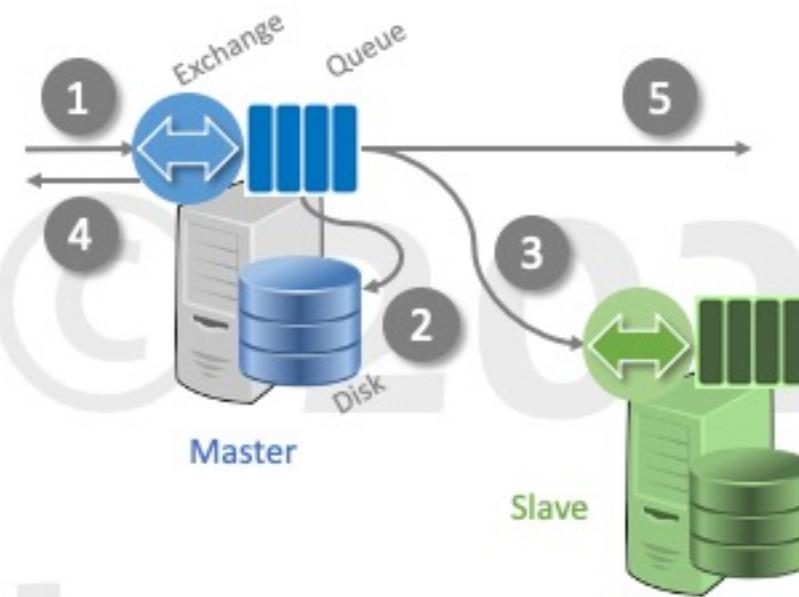


- Design Goals
  - At-Least-Once Delivery
  - Message Sequencing – FIFO
  - Interface Decoupling
  - Consumer Decoupling
  - Message Rate Decoupling
- Consumer Mode
  - Push
  - Pull
- Use Cases
  - Service Integration
  - Message Buffer



# Rabbit MQ

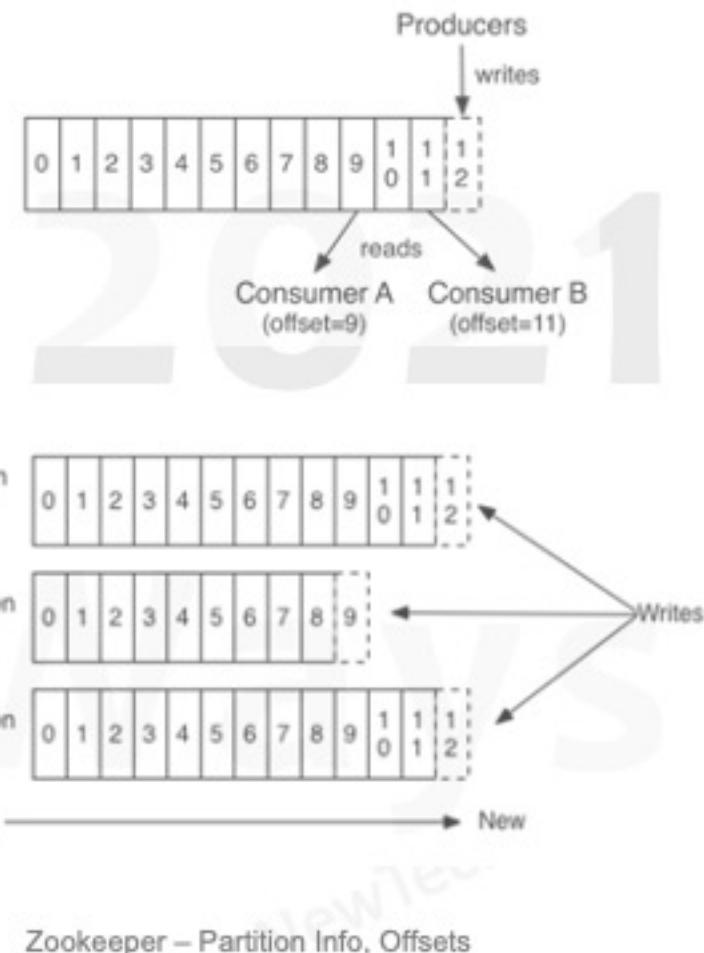
- General purpose message broker
- Messages are pushed to consumers
- Message deleted once acknowledged
- Message ordering is guaranteed
- Useful for asynchronous service integration
- Both persistent and transient messages are supported
- Uses built-in component called exchange for routing



- Scales to 50K messages per second
- Scales well vertically
- Not horizontally scalable
  - Master-Slave replication for high availability
    - Clients can connect to any node
    - All publish, consume operations first go to master

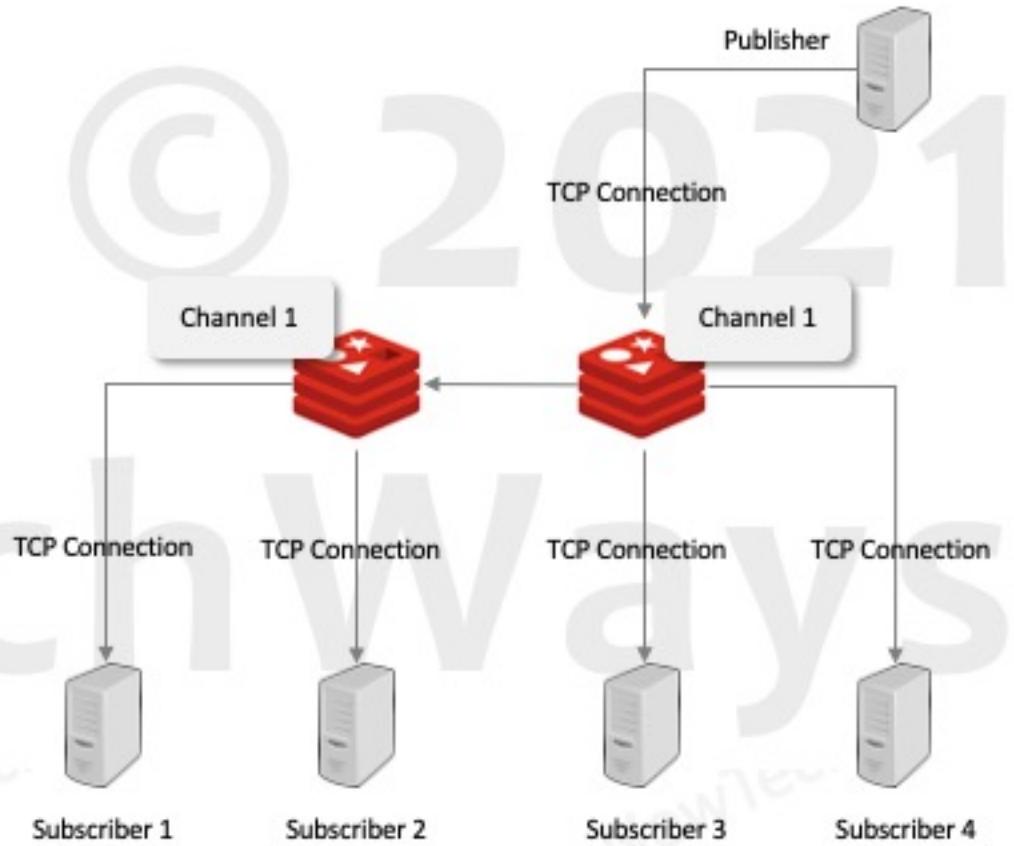
# Kafka

- Performance
  - Million messages per second
  - Sequential writes and reads on log files
  - Relies on page cache for quick reads
- Horizontally Scalable
  - Topics are partitioned
- Order of data guaranteed only within a partition
  - Producer can decide the partition
- Messages are not deleted, so can be replayed
- Consumers can only pull the data
  - No push by Kafka
  - Not designed for service integration
- Useful for streaming analytical workloads
  - Where high throughput is required
  - Click streams, Page views, Logging, Ingestion, Security



# Redis Pub/Sub

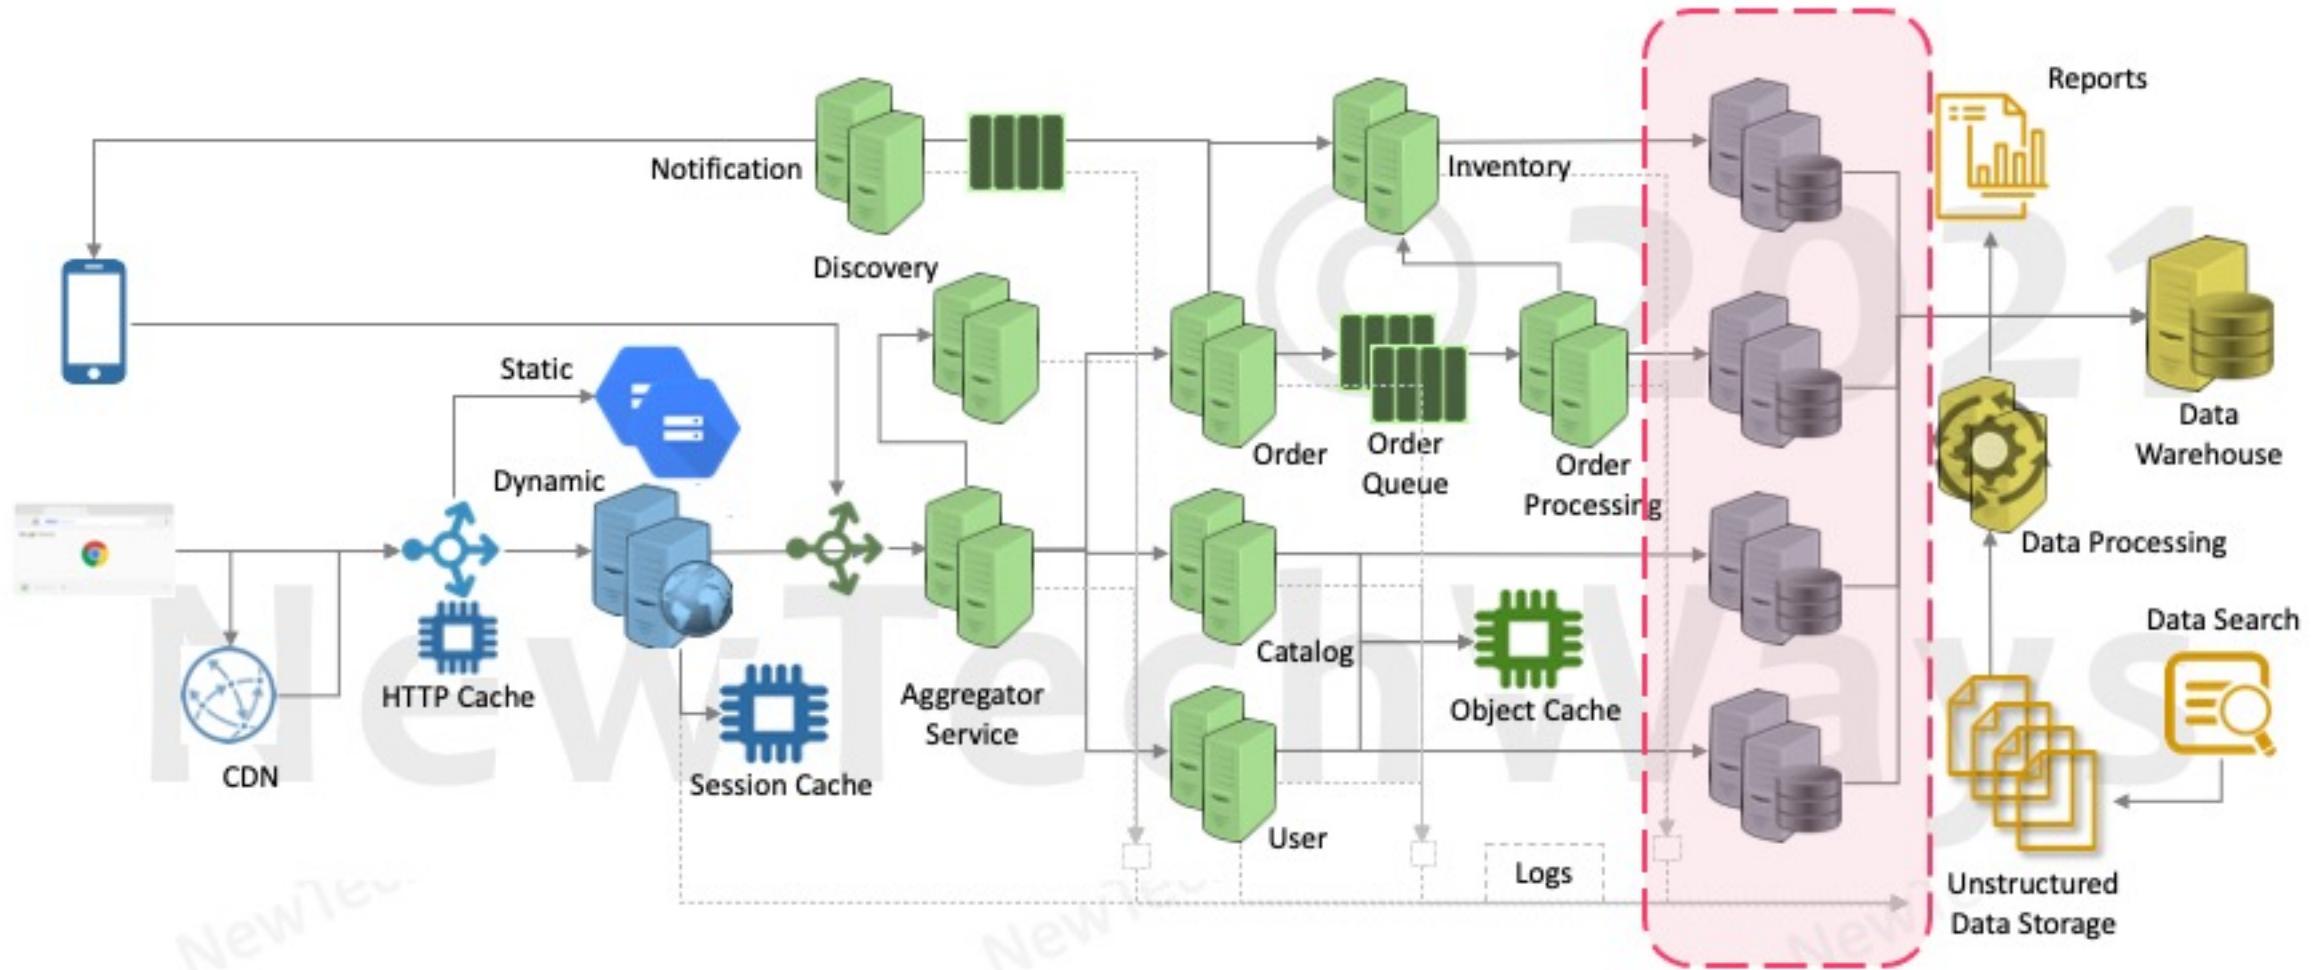
- For short lived messages with no persistence
  - Much like a synchronous call
  - Fire and forget
    - No delivery guarantee
- Million operations per second
- Useful for making dashboards
  - Leaderboard
- Comparison
  - Kafka
    - No push
    - Writes to log
  - RabbitMQ Transient
    - Delivery acknowledgement
    - Deletion of delivered messages



# Cloud MQ Solutions

| Community | AWS     | Google Cloud |
|-----------|---------|--------------|
| Rabbit MQ | SQS     | Pub/Sub      |
| Kafka     | Kinesis |              |

# Datastores



# Datastore – Solutions

- RDBMS
  - Oracle
  - SQL Server
- Distributed Databases
  - Key-Value
    - Dynamo
  - Column Family
    - Big Table
    - Cassandra
    - HBase
  - Document Oriented
    - MongoDB

© 2021

NewTechWays

# RDBMS

- General purpose database (< 1-5 TB data, 10K connections)

- ACID Transactions

- Update of multiple records or tables

Possible only on a single node or else it requires 2PC/3PC

- Data Consistency

- Data same for all readers at any given time

Leads to low availability

- Fixed Schema

- Data analysis

- Columns can be selected
    - Rows can be filtered
    - Queries can Join Tables

Impedes application evolution

- Query pattern can change or evolve

- Any column can be indexed
    - Indexes can be created whenever required

Joins may slow down a system

- Normalized Data

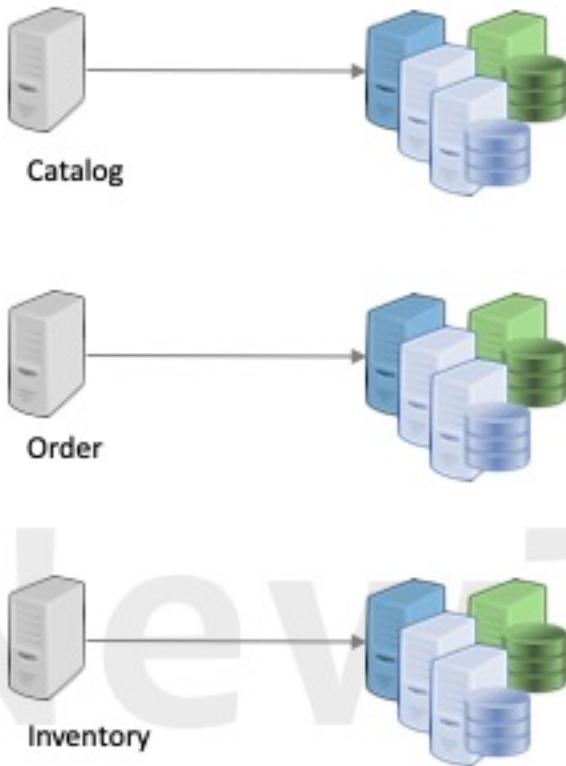
- Efficient storage and writes

Allows only one version of data

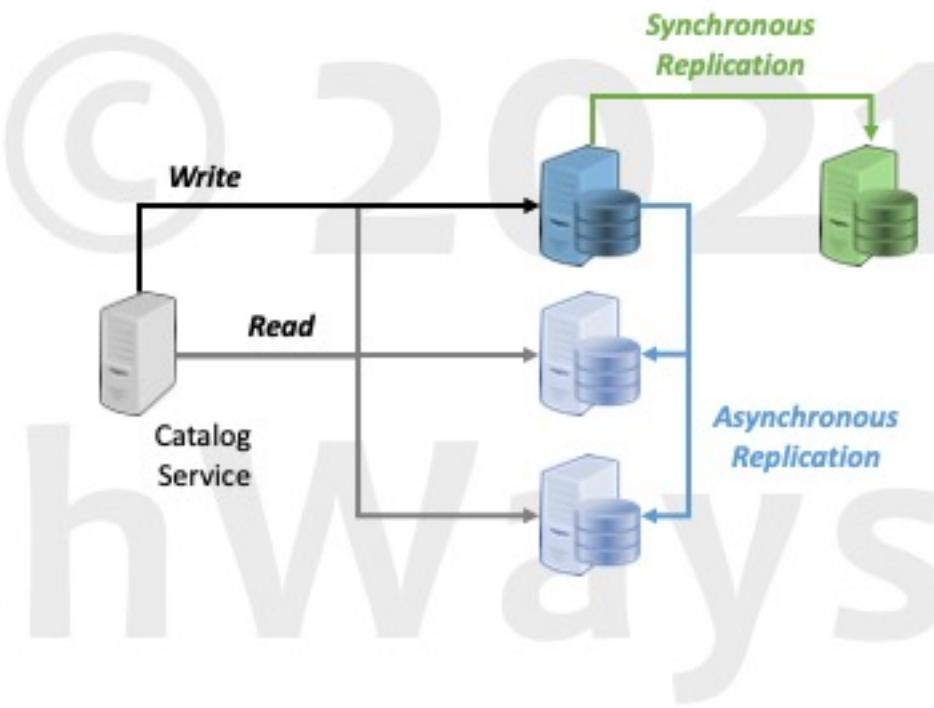
- Overwrite for updates

Old design for expensive disk space

# RDBMS Scalability Architecture



Vertical Partitioning for Scalability



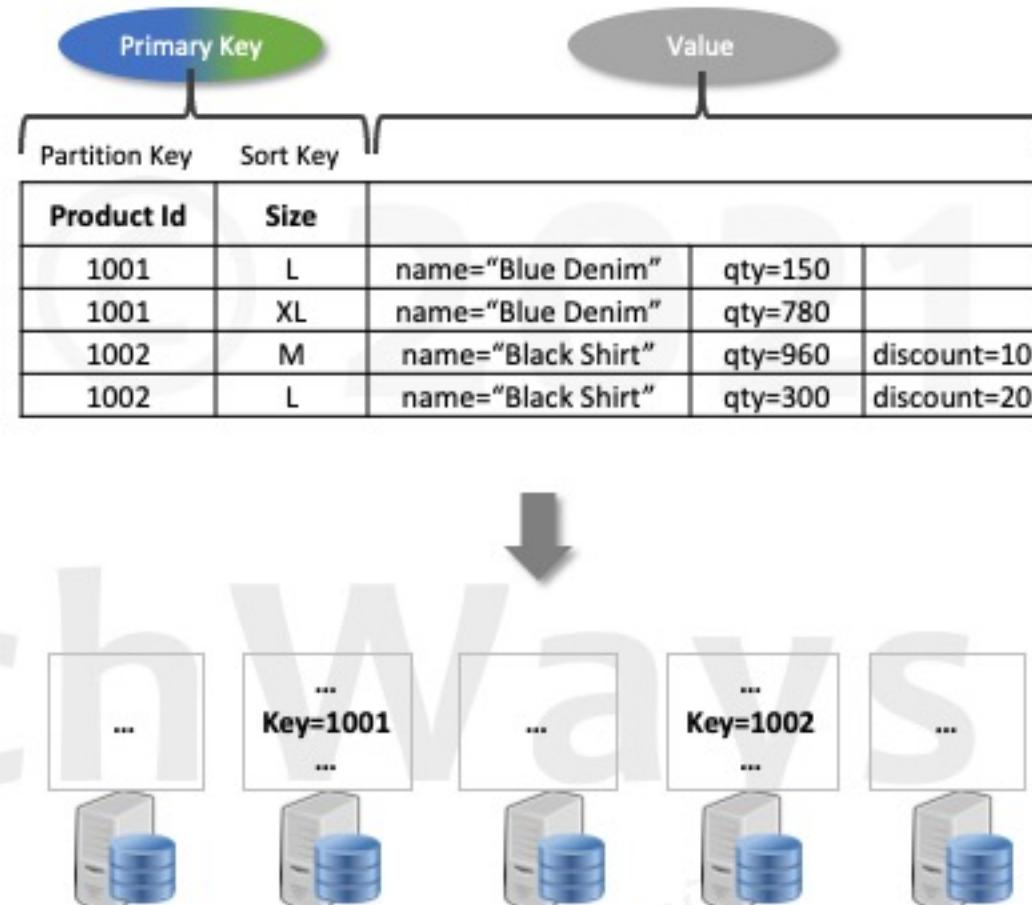
Replication for Reliability/Availability

# NoSQL Objectives & Trade-Offs

| Objective       | Architectural Choice                           | Trade-Off                                        |
|-----------------|------------------------------------------------|--------------------------------------------------|
| Scalability     | Horizontal Partitioning,<br>Commodity Hardware | ACID Transactions,<br>Joins                      |
| Availability    | Data Replication,<br>Eventual Consistency      | Data Consistency                                 |
| Flexible Schema | Key-Value, Column Family,<br>Document-Oriented | SQL, Secondary Indexes,<br>Integrity Constraints |
| Performance     | Aggregate Schema,<br>In-Memory R/W             | Normalization,<br>Non-Key Queries                |

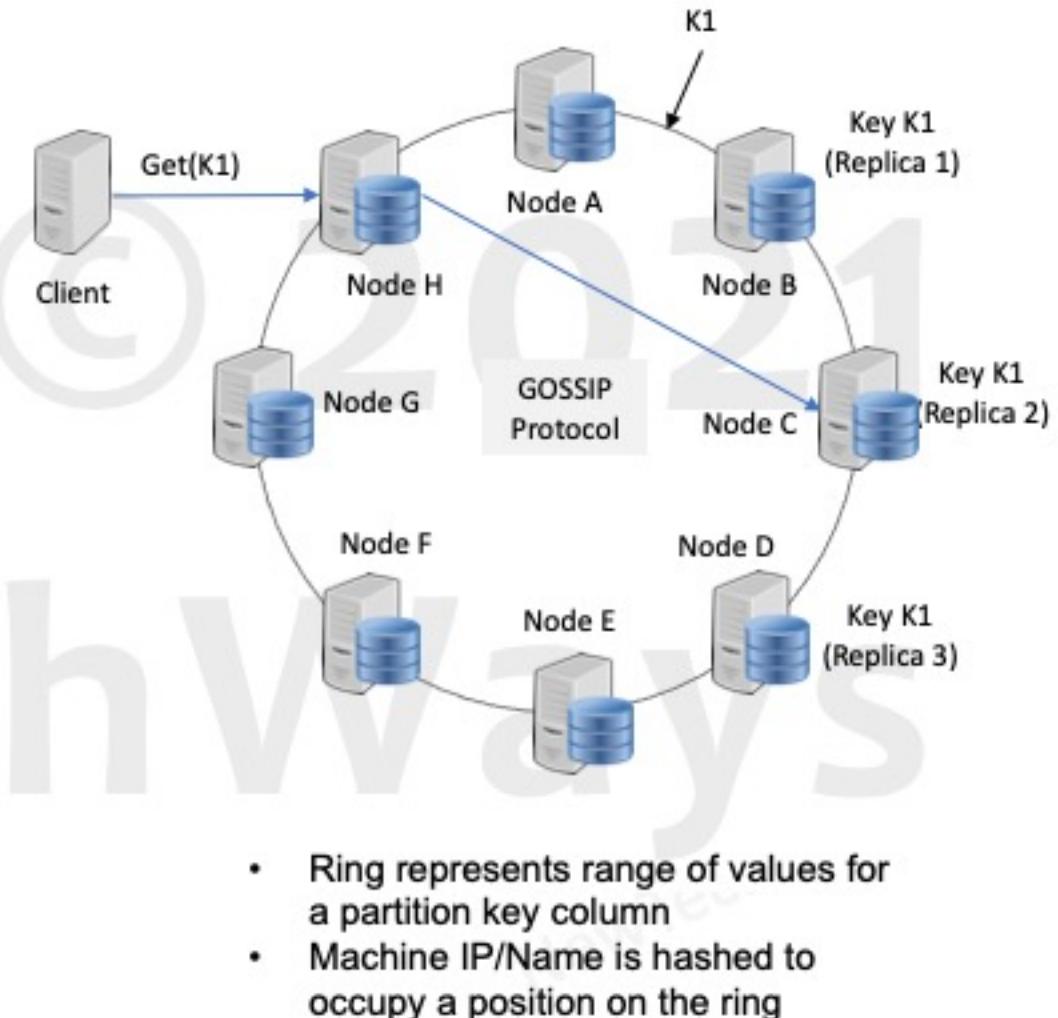
# Amazon DynamoDB

- Key-Value Pair Datastore
- Used for High Scalability & Availability
- Table is a Hash-Map
  - Persistent
  - Distributed
- API
  - Put, Get, Update, Delete, Query
- Index Key – has two parts
  - Partition Key
    - Can have only one attribute
    - Key is hashed to determine the partition
  - Sort Key
    - Determines sort order of an item within a partition
    - Can be used for range query <, >, like
- R/W ops for a key are atomic



# Dynamo DB Architecture

- Suitable for storing small chunks of data
  - Key values less than 1MB
- Peer-To-Peer cluster
  - No master – write on any node
  - Each node responsible for a set of keys
  - Consistent hashing of virtual nodes for key allocation
- Highly Scalable
  - Practically any number of nodes
  - Petabytes of data
  - Can handle 10 million RW ops requests/second
- Highly Available
  - Updates are not rejected even in case of network partitions or server failures
  - Vector clocks and business rules to resolve merge conflicts
- Consistency guarantee is adjustable
  - $R+W > N$  for strong consistency



# Google Big Table

- Basis for Apache HBase
- Column-Family Storage
- Table as Tree-Map
  - Sparse
  - **Sorted**
  - Persistent
  - Distributed
- Limited CF (< 100)
- CFs are compressed
- Unlimited columns
- R/W ops atomic for a key
- Timestamps for versioning

| Row Key          | Document:                   | Language: | Referrer: trips.com | Referrer: holidays.com | Referrer: news.com |
|------------------|-----------------------------|-----------|---------------------|------------------------|--------------------|
| com.airlines.www | <!DOCTYPE html><br><html>.. | EN        | Check Flights       |                        |                    |
| com.hotel.www    | <!DOCTYPE html><br><html>.. | EN        |                     | Your Hotel             | Venue              |

Index Key: Row Key/Col Family:Col Name/timestamp

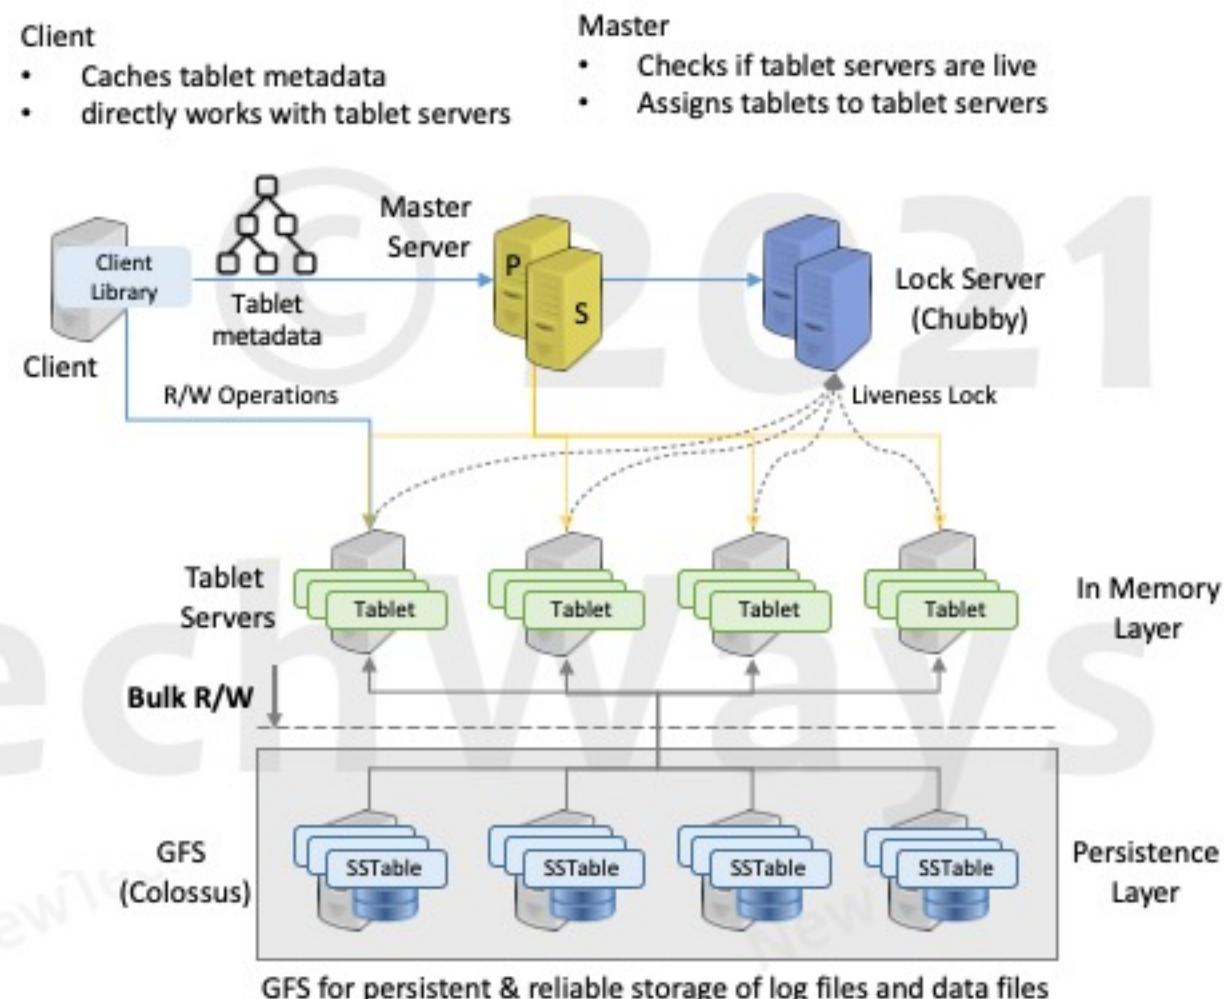
| Row Key          | Column Family | Column       | Timestamp | Value                   |
|------------------|---------------|--------------|-----------|-------------------------|
| com.airlines.www | Document      |              | 987654321 | <!DOCTYPE html><html>.. |
| com.airlines.www | Language      |              | 987654321 | EN                      |
| com.airlines.www | Referrer      | trips.com    | 987654321 | Check Flights           |
| com.hotel.www    | Document      |              | 987655432 | <!DOCTYPE html><html>.. |
| com.hotel.www    | Language      |              | 987655432 | EN                      |
| com.hotel.www    | Referrer      | holidays.com | 987655432 | Your Hotel              |
| com.hotel.www    | Referrer      | news.com     | 987655432 | Venue                   |

# Bigtable Architecture

- Schema-less, structured (CF), sorted data
- GFS for reliable persistent storage
  - Replicated storage for large files
- In-memory tablet servers
  - High R/W throughput with low latency
  - Single copy of data for read and write
    - No write conflicts
- Horizontally scalable
  - Stores peta/exabytes of data
  - Client load is distributed
- Strongly consistent

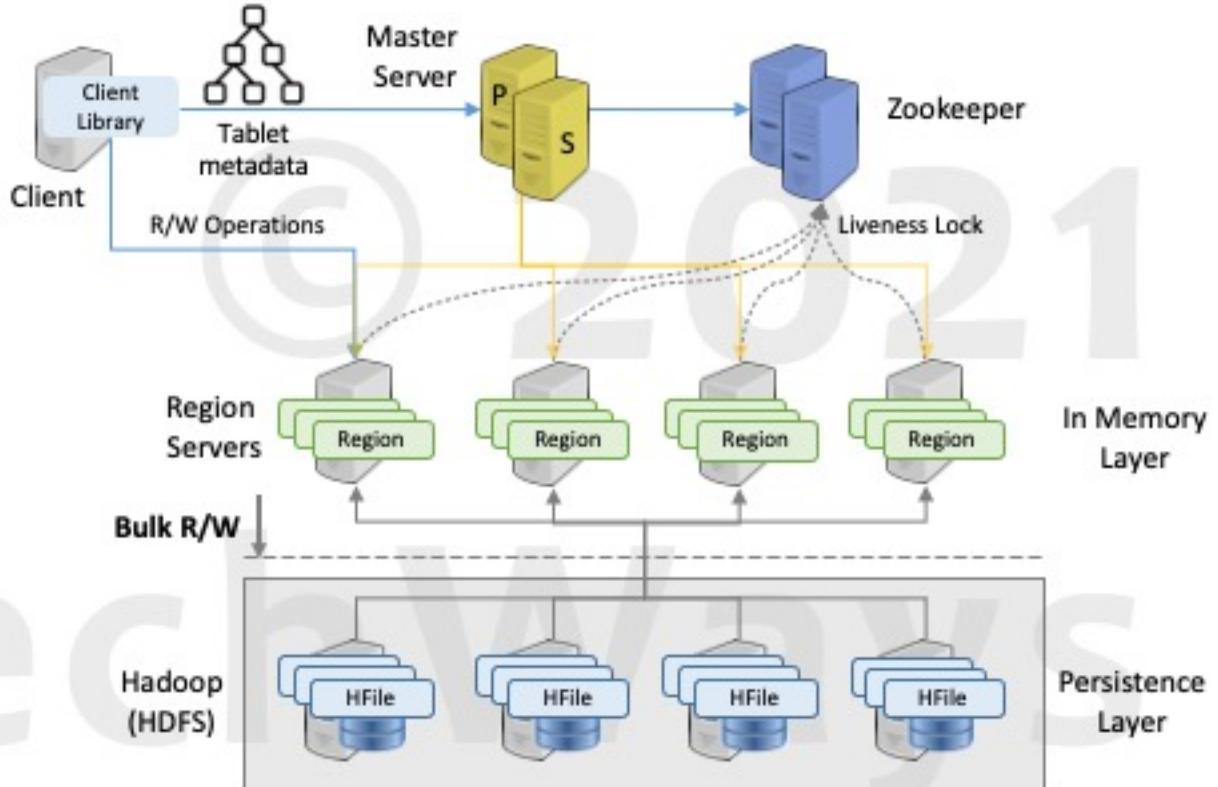
| New Key          | Column Family | Column     | Type   | Value                      |
|------------------|---------------|------------|--------|----------------------------|
| com.airlines.www | Document      |            | String | 1000+ flight destinations. |
| com.airlines.www | Language      |            | String | EN                         |
| com.airlines.www | Referrer      | trips.com  | String | Check Flights              |
| com.hotel.www    | Document      |            | String | 1000+ travel websites.     |
| com.hotel.www    | Language      |            | String | EN                         |
| com.hotel.www    | Referrer      | hotels.com | String | New Hotel                  |
| com.hotel.www    | Referrer      | news.com   | String | Vacation                   |
| com.airlines.www | Document      |            | String | 1000+ flight destinations. |
| com.airlines.www | Language      |            | String | EN                         |
| com.airlines.www | Referrer      | trips.com  | String | Check Flights              |
| com.hotel.www    | Document      |            | String | 1000+ travel websites.     |
| com.hotel.www    | Language      |            | String | EN                         |
| com.hotel.www    | Referrer      | hotels.com | String | New Hotel                  |
| com.hotel.www    | Referrer      | news.com   | String | Vacation                   |

Data ranges are split into tablets



# HBase

- Open source implementation of Bigtable
- Column Family schema
- Keys are Range Partitioned
- Storage on Hadoop HDFS
- Strong consistency over high availability
- Highly scalable
- High throughput and low latency writes



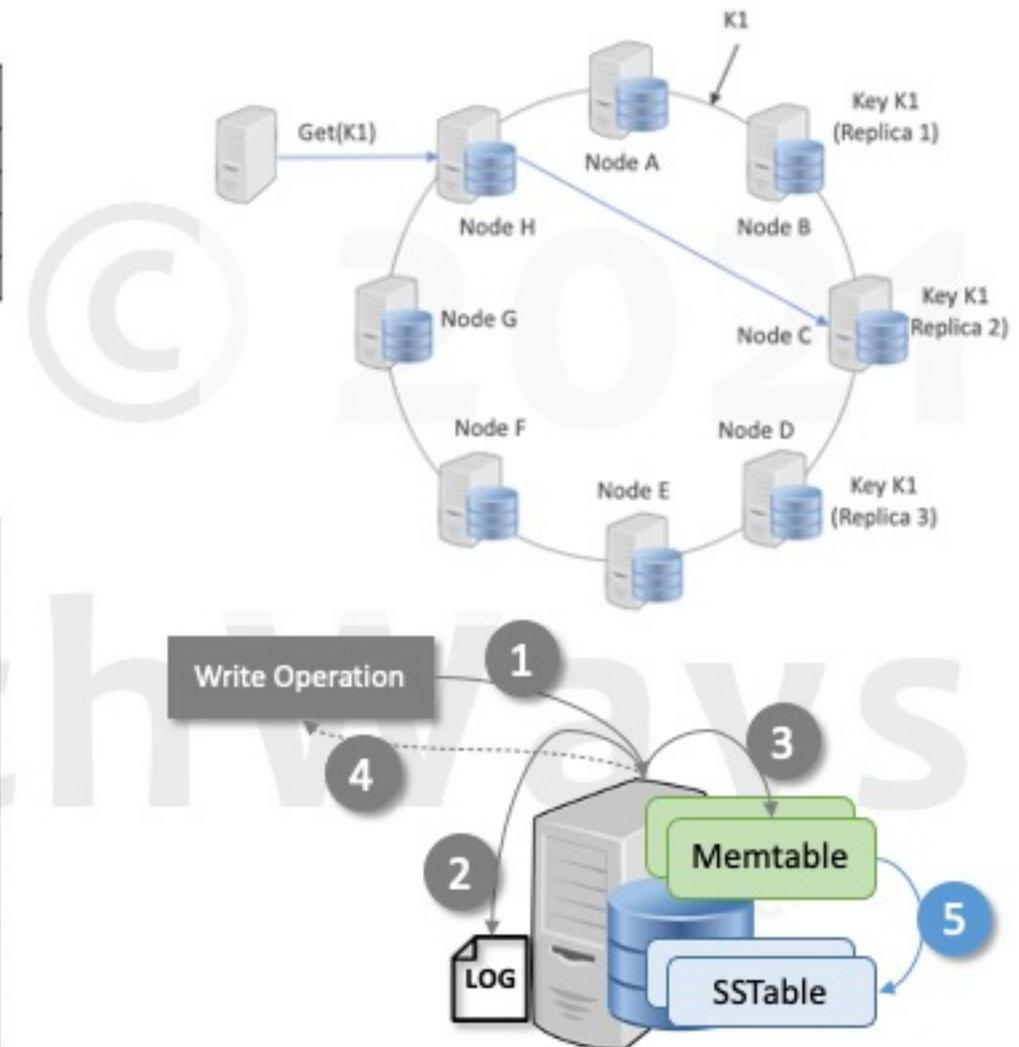
# Cassandra

| Partition Key | Sort Key | Columns     |          |          |
|---------------|----------|-------------|----------|----------|
| Product Id    | Size     | Name        | Quantity | Discount |
| 1001          | L        | Blue Denim  | 150      |          |
| 1001          | XL       | Blue Denim  | 780      |          |
| 1002          | M        | Black Shirt | 960      | 10       |
| 1002          | L        | Black Shirt | 300      | 20       |

Column Family

| Product Id | Size | Col Name | Col Value   |
|------------|------|----------|-------------|
| 1001       | L    | Name     | Blue Denim  |
|            | L    | Quantity | 150         |
| 1001       | XL   | Name     | Blue Denim  |
|            | XL   | Quantity | 780         |
| 1002       | M    | Name     | Black Shirt |
|            | M    | Quantity | 960         |
|            | M    | Discount | 10          |
|            | L    | Name     | Black Shirt |
| 1002       | L    | Quantity | 300         |
|            | L    | Discount | 20          |

Partition +  
Partition -



# Cassandra Features

- Schema-less, Column-Family Structured Data

- Sparse
- Persistent
- Distributed

- Horizontally Scalable

- Petabytes of data

- Highly Available

- Even during network partitions

- High throughput R/W operations

- Merge conflicts

- Vector clocks
- Timestamps
- Business rules

| Data Model        | Column Family         | Big Table |
|-------------------|-----------------------|-----------|
| Storage           | Memtable → SSTable    | Big Table |
| Partitioning      | Hash Based            | Dynamo    |
| Replication       | Peer-to-Peer          | Dynamo    |
| Consistency Model | Eventually Consistent | Dynamo    |
| Availability      | Highly Available      | Dynamo    |

# MongoDB

- Key -> Document
  - In Binary JSON (bson) format
- Columns created dynamically
- Columns can be indexed
- Documents can be queried
  - on id and/or column values
- An operation on a single document is atomic
  - 2 PC for multiple documents

```
db.inventory.insertMany([
 { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },
 { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },
 { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }
]);

db.inventory.find({ status: "A", qty: { $lt: 30 } }) Insert Query

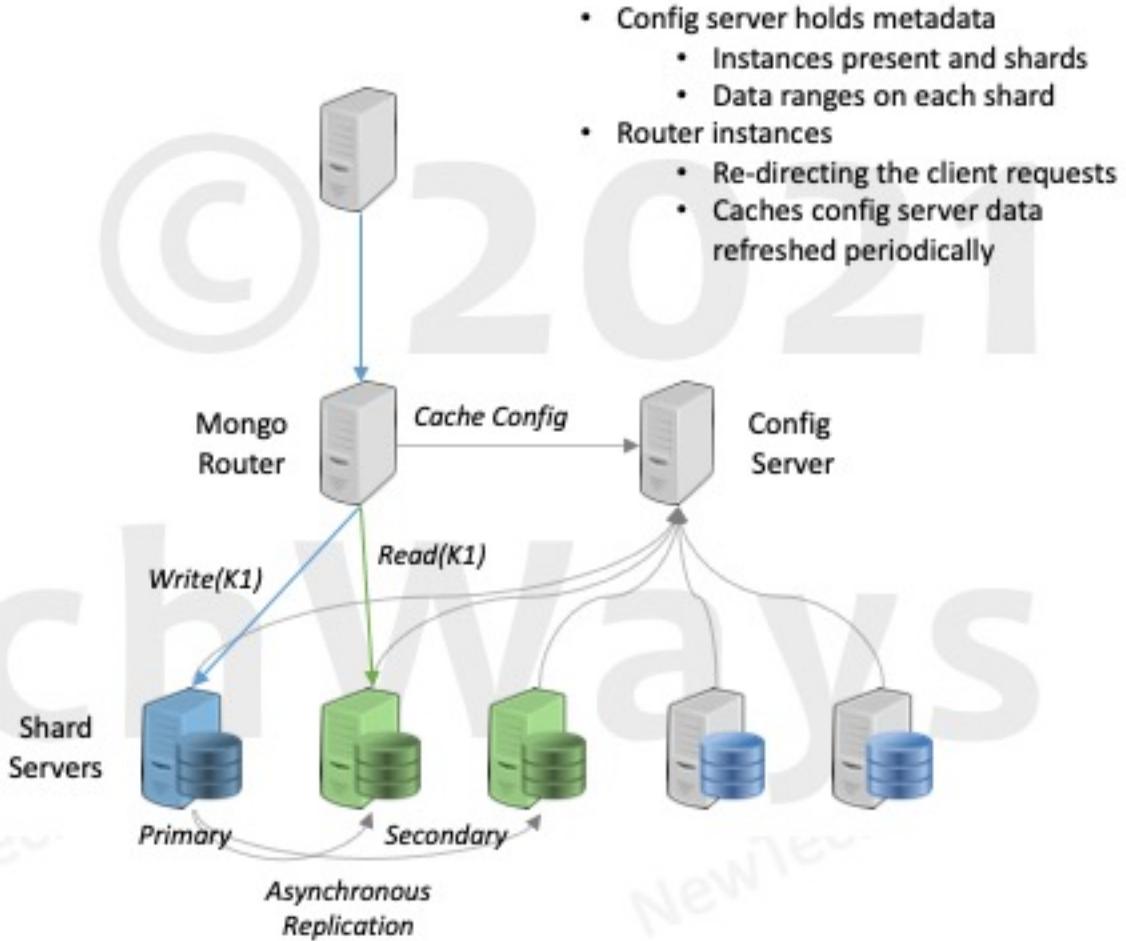
db.inventory.find({ "size.h": { $lt: 15 } }) Query with nested attribute

db.inventory.updateOne(
 { item: "paper" },
 {
 $set: { "size.uom": "cm", status: "P" },
 $currentDate: { lastModified: true }
 }
) Update

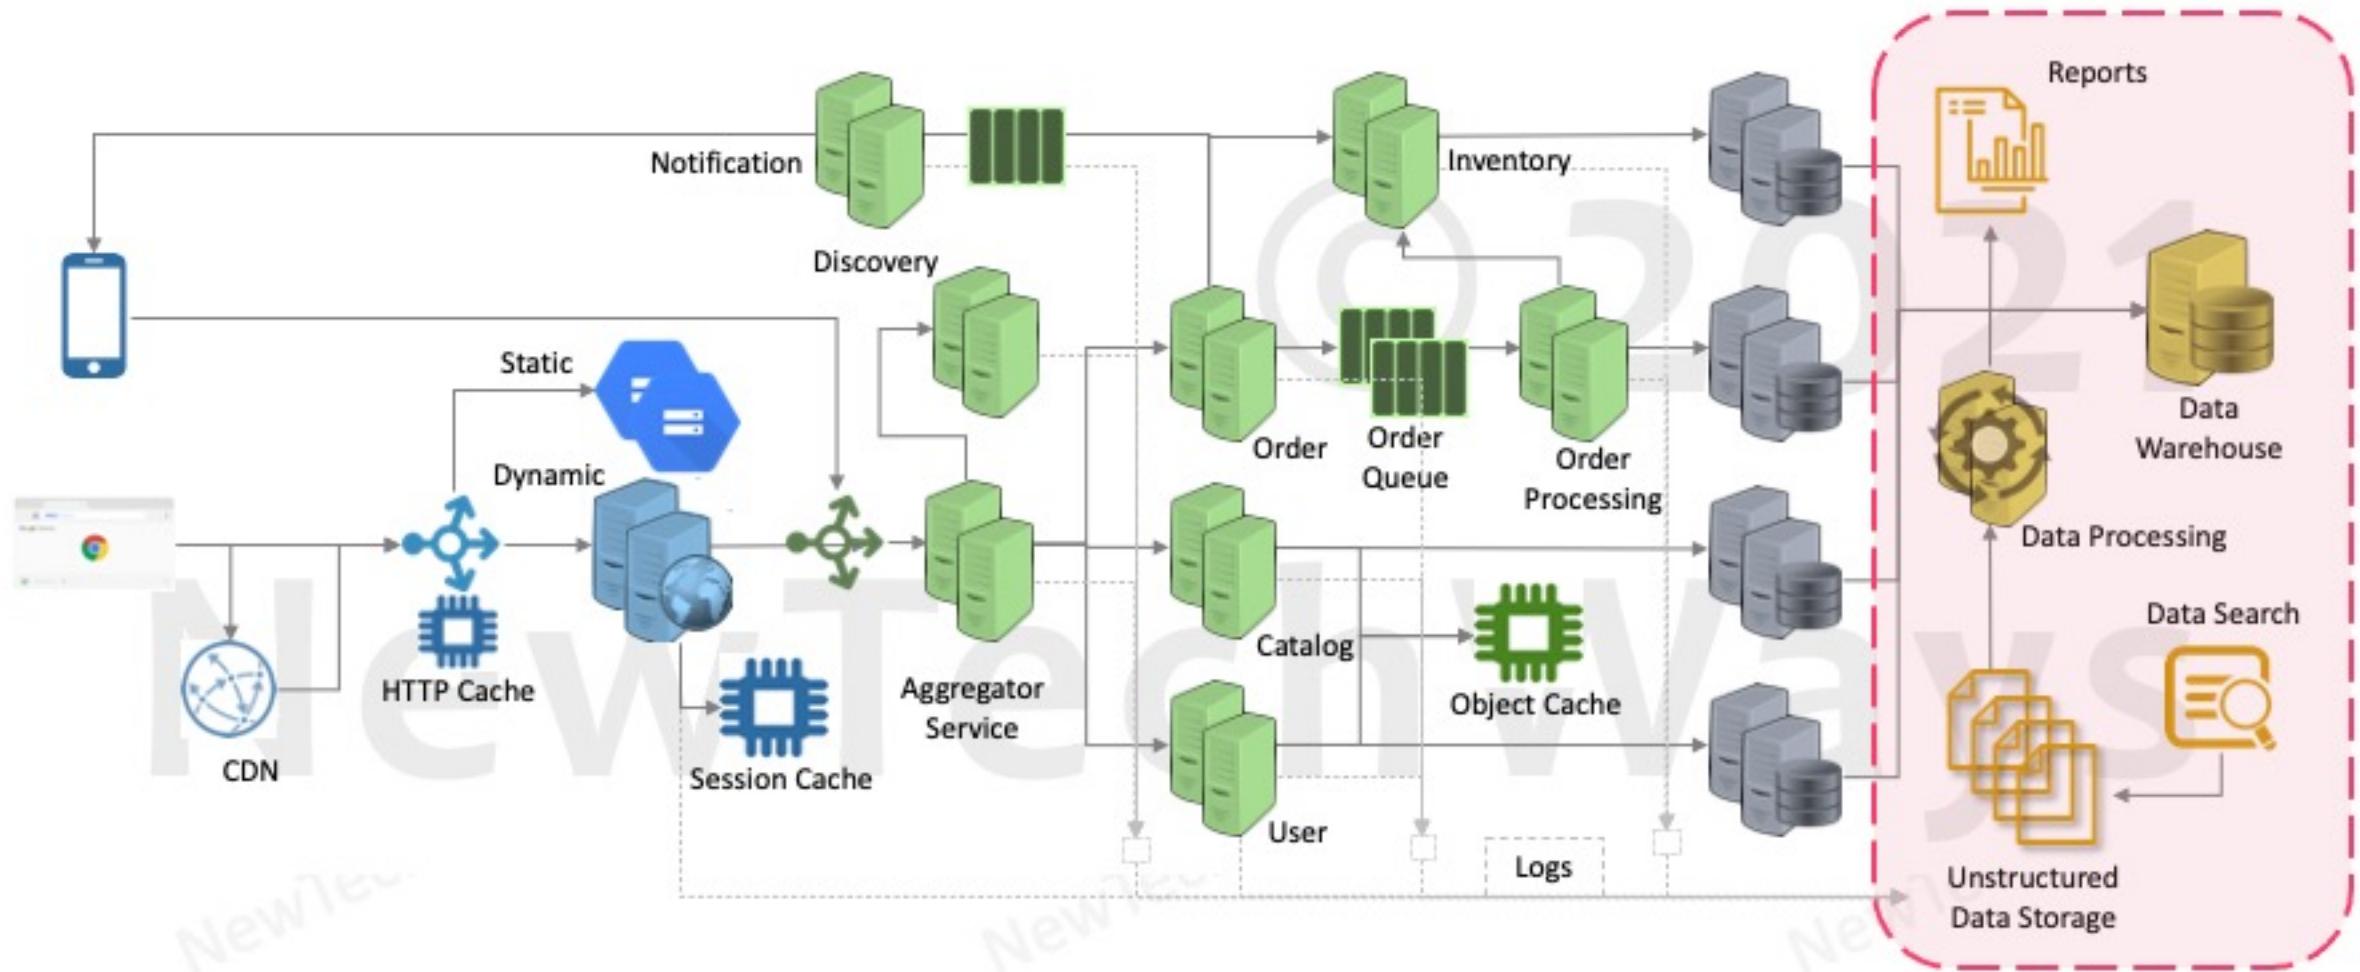
db.products.createIndex(
 { item: 1, quantity: -1 } Index Creation
)
```

# MongoDB Architecture

- Indexing for fast search
  - It's a write overhead
- Sharding for Scalability
  - Range sharding
  - Hash sharding
- Replication
  - Master Slave
    - No write conflicts
  - Asynchronous (default)
    - Eventual consistency
  - Synchronous (on demand)
- Works very well with Node.js
  - Javascript -> Node.js -> Mongo
    - JSON format



# Analytics



# Analytics – Solutions

- Data Movement
  - Logstash
  - Fluentd
- Storage
  - Hadoop HDFS
    - Map-Reduce
  - Elastic Search
    - Search
- Stream Processing
  - Kafka
    - Buffer
  - Storm, Flink
    - Processing

© 2021

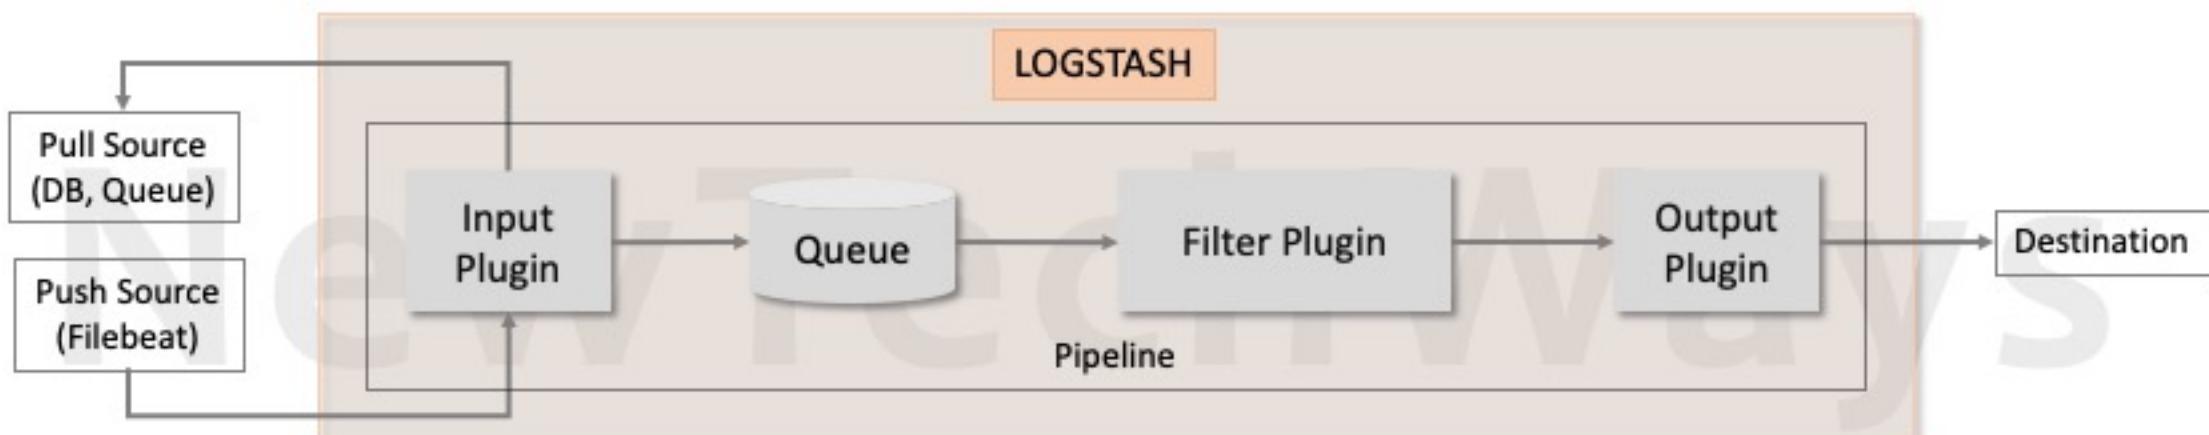
NewTechWays

# Logstash Architecture

- I/O Plugins
  - Collect & Move log/stream events
- Queue
  - Reliable data movement
  - At least once delivery
  - Tracks Acks, Applies Backpressure
- Filter Plugins
  - Filter, Transform, Aggregate

127.0.0.1 - - [05/Feb/2012:17:11:55 +0000]  
"GET / HTTP/1.1" 200 140 "-" "Mozilla/5.0..."

{"  
"host": "127.0.0.1",  
"user": "-",  
"method": "GET",  
"agent": "Mozilla/5.0 (Windows..."}



Access Logs  
App Logs  
System Logs  
DB Logs ...

File, Beats,  
Http, JMS,  
Redis, Kafka,  
ES, ...

Memory  
File

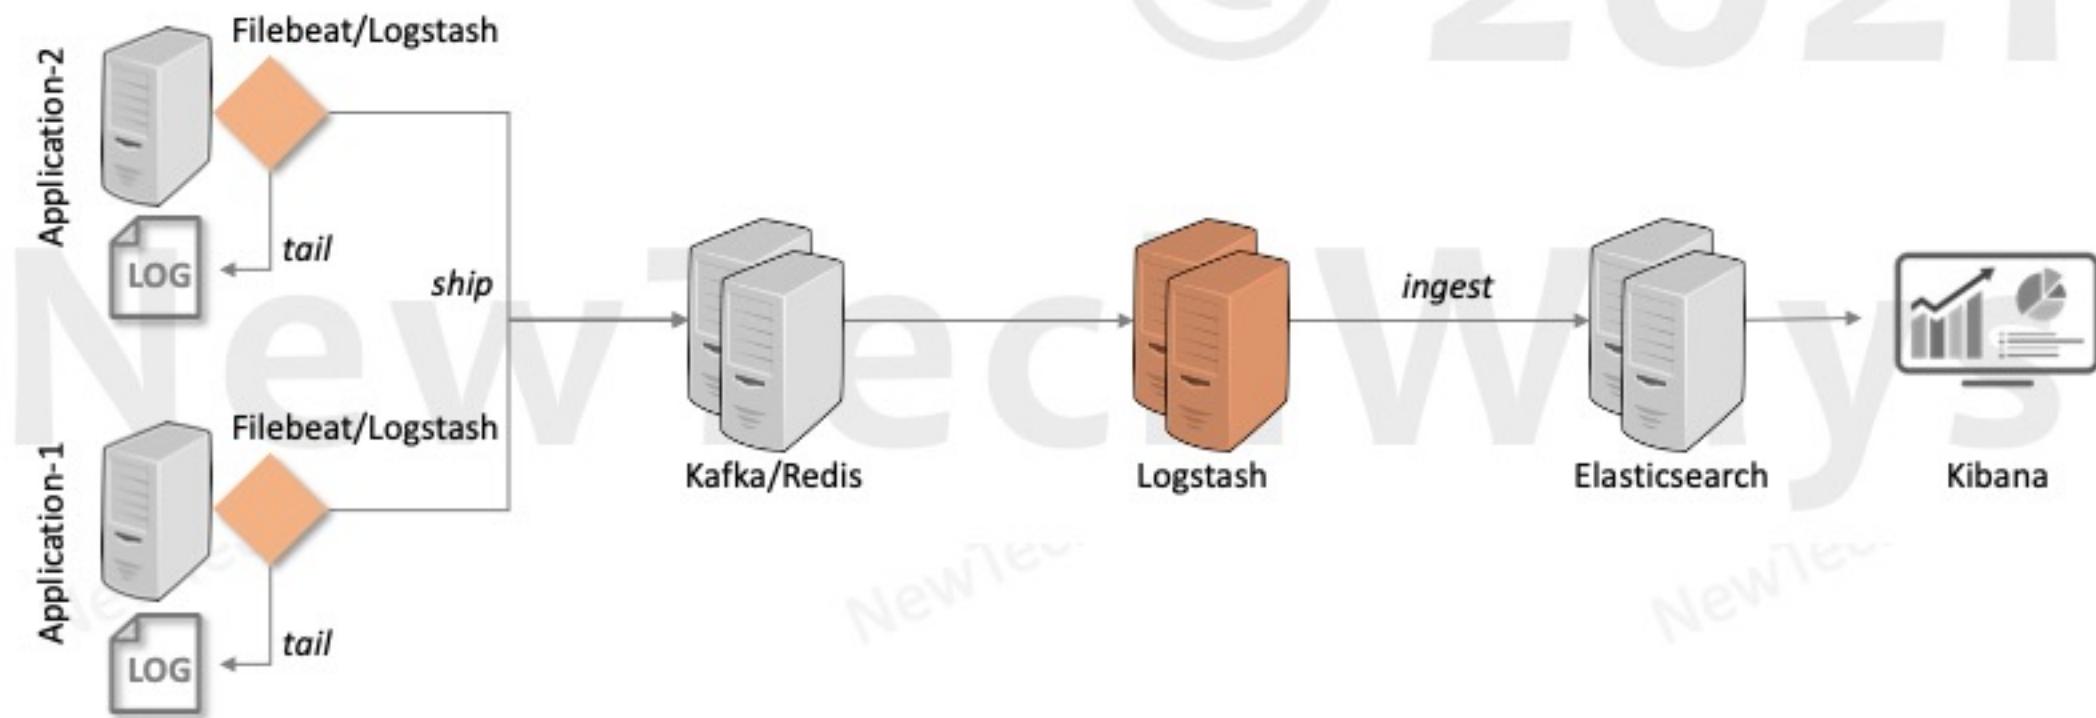
JSON, CSV, Grok,  
Dissect, Clone, Drop,  
Aggregate, ...

ES, Redis,  
Webhdfs, S3,  
Kafka, File,  
Email, ...

Alerts  
Analysis  
Archives

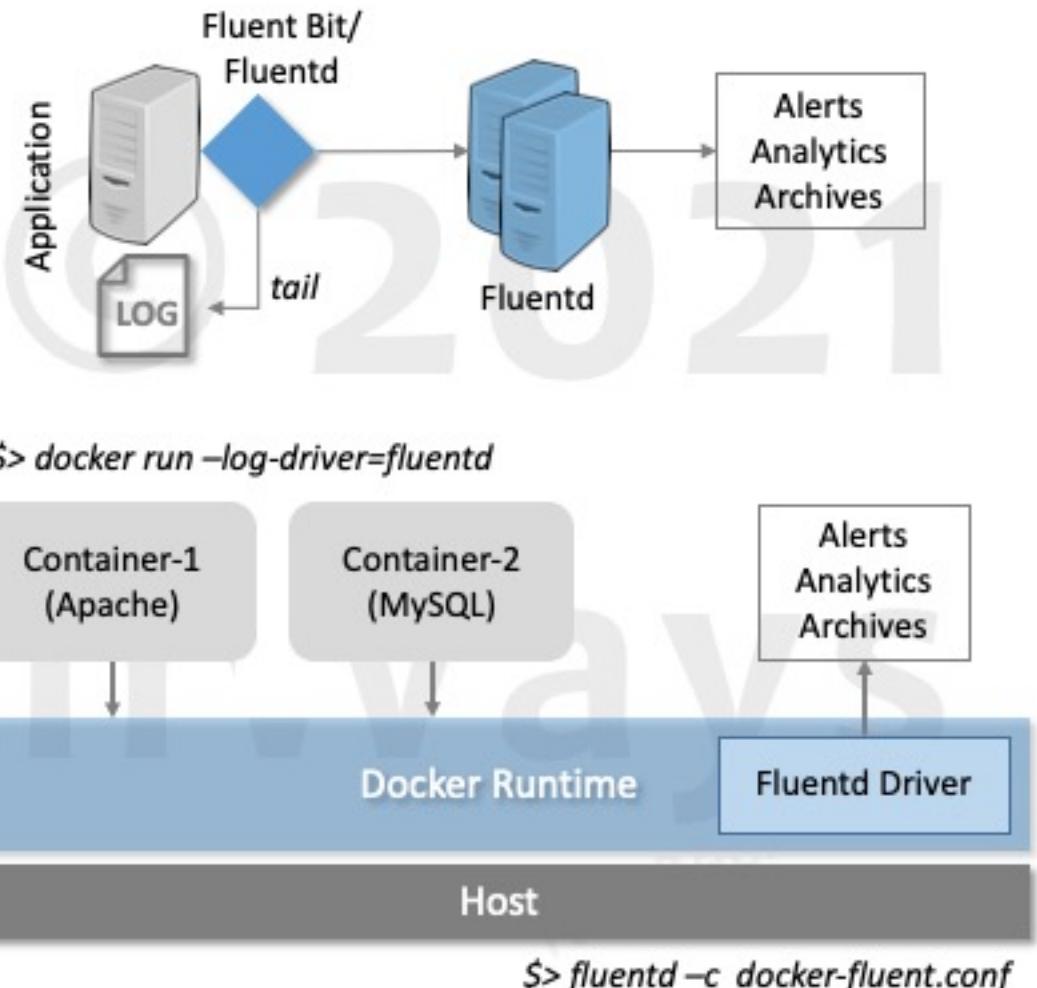
# Logstash Data Streaming Architecture

- Streaming log data for Real-Time Analytics
- Horizontally Scalable & Highly Available – Any number of Logstash nodes
- Fault Tolerance – Needs reliable disk storage (RAID, Cloud persistent disks)
- Use Kafka as buffer – For heavy load that Logstash Queue cannot handle



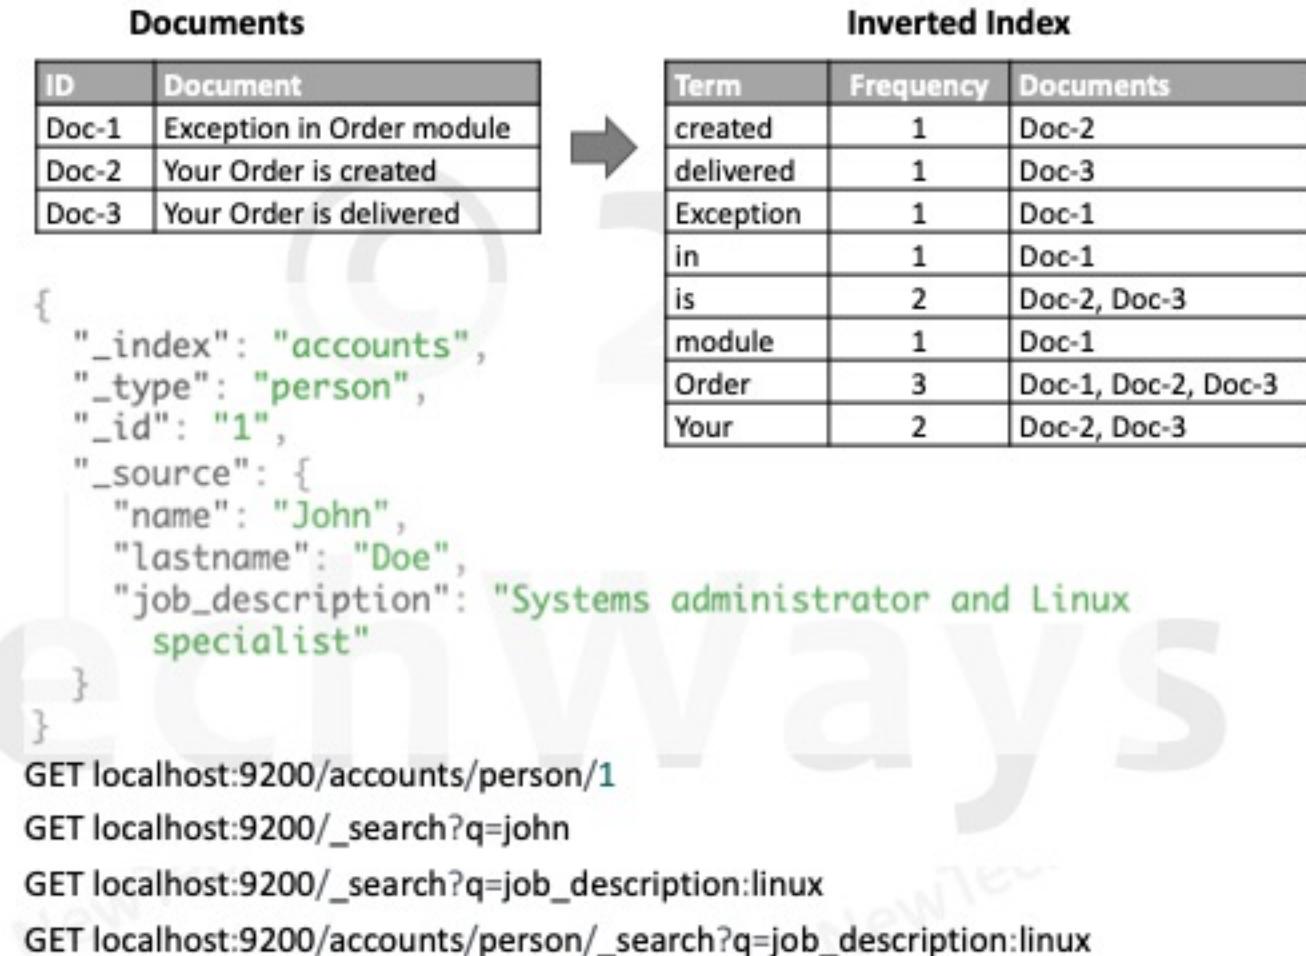
# Fluentd

- Older than Logstash
- Memory Footprint
  - Logstash – Heavyweight (GB)
  - Filebeat – Lightweight (MB)
  - Fluentd – Lightweight (MB)
  - Fluent Bit – Super Lightweight (KB)
- All features of Logstash
- + Routing
  - Tags
- + Docker Logging
  - Picks log events from container console



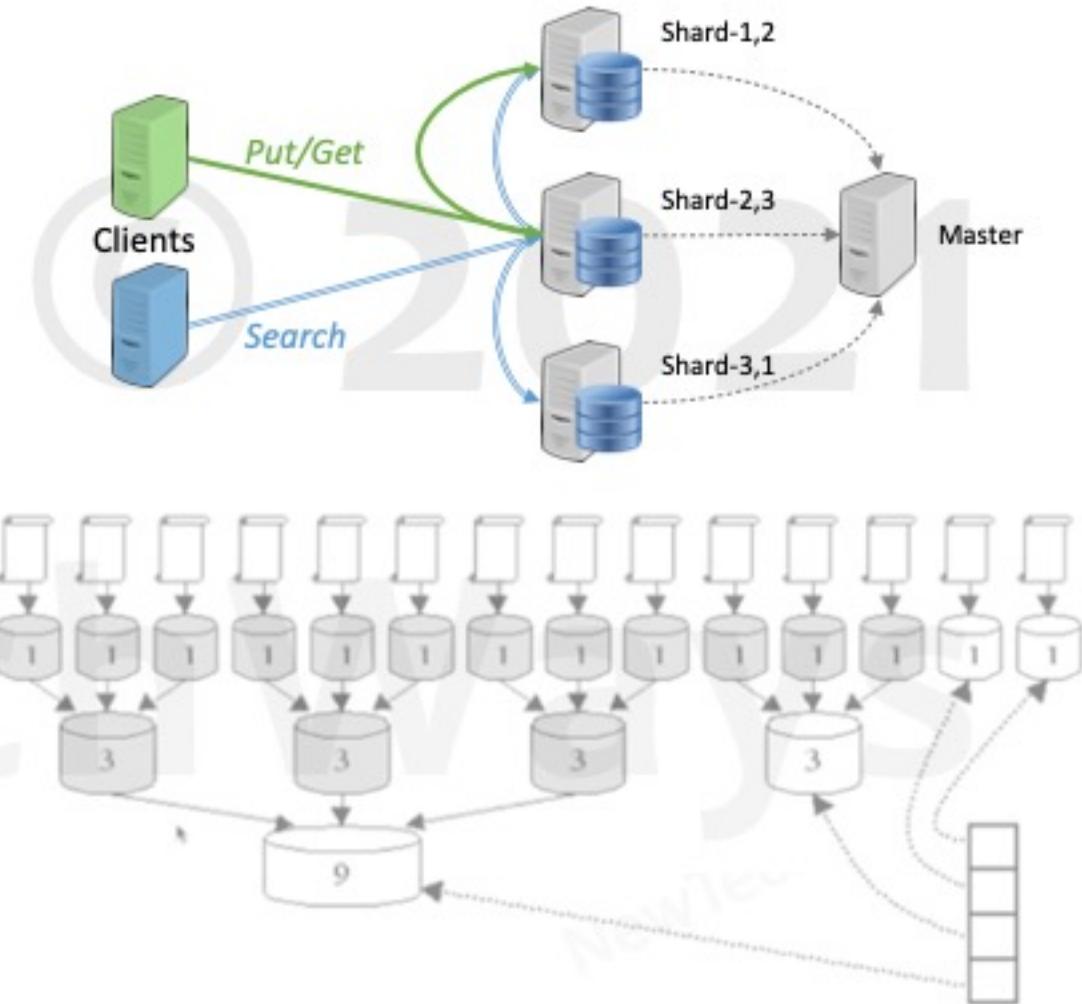
# Elasticsearch

- Full-Text Search
  - Filter, Group, Aggregate
- Stores JSON Documents
- Document Fetched using id
- Indexes JSON keys and values
- Structure
  - Index -> Database
  - Type -> Table
  - Document -> Row
    - JSON keys are flattened
  - Supports data types
- Users can specify mapping between terms & documents



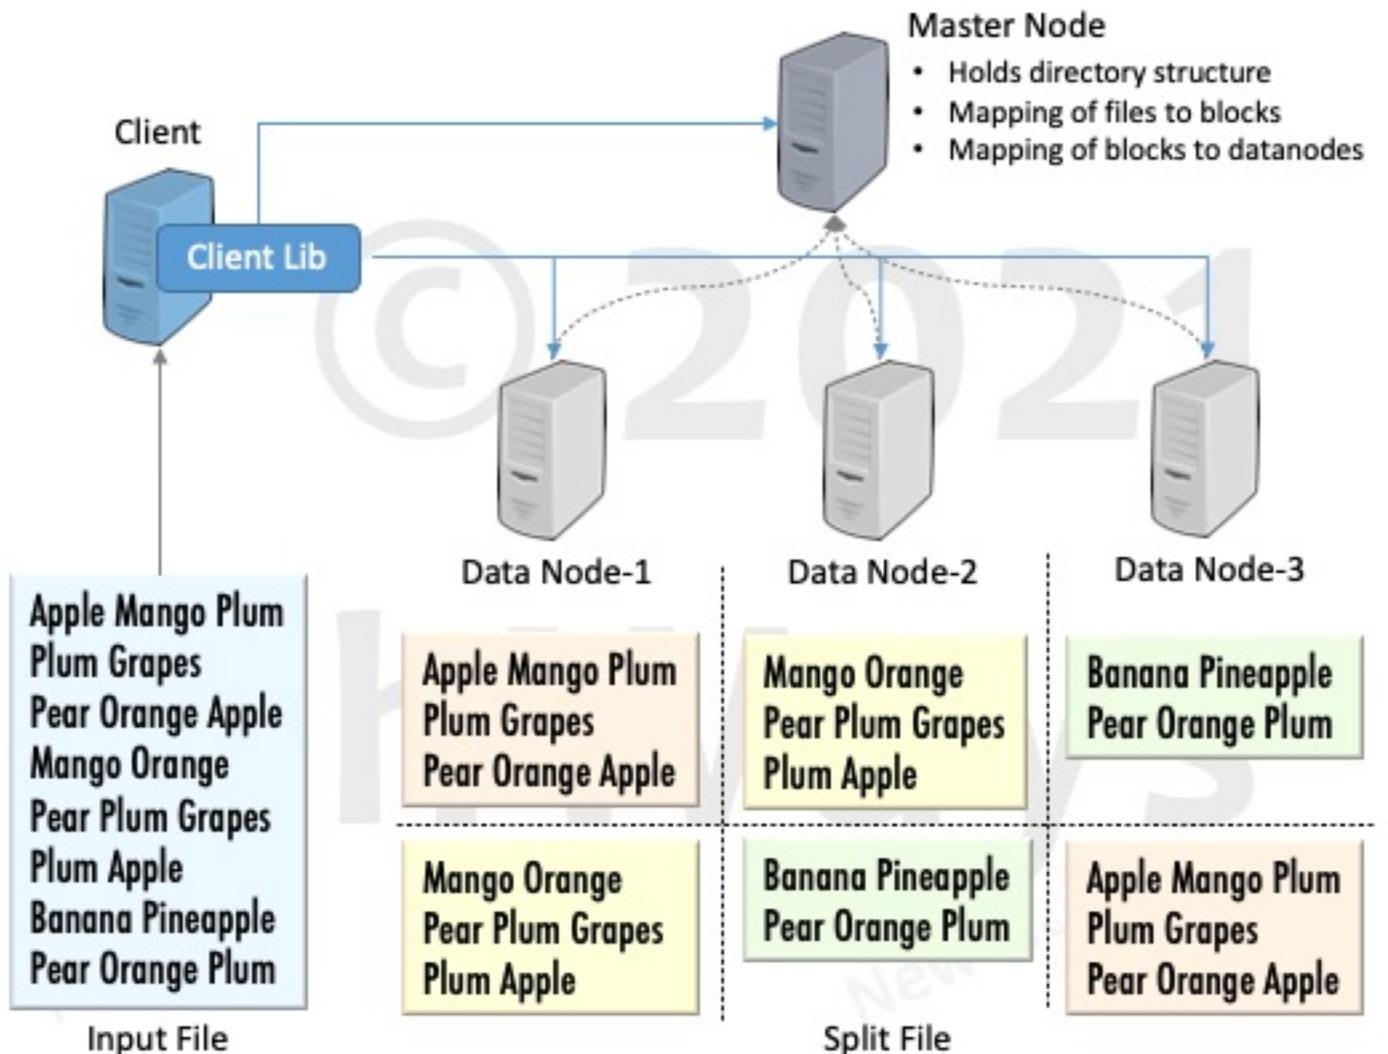
# Elasticsearch Architecture

- Document Oriented data-model
- Horizontally Scalable
  - Petabytes of data
  - Data is sharded with key as Document Id
  - Put/Get request goes to specific shard
  - Search queries go to all shards
- Highly Available
  - Shards are replicated
- Index structure is based on merge sort
- Index not updated with every update or insert
- Index maintained in memory
  - Occasionally flushed to disk



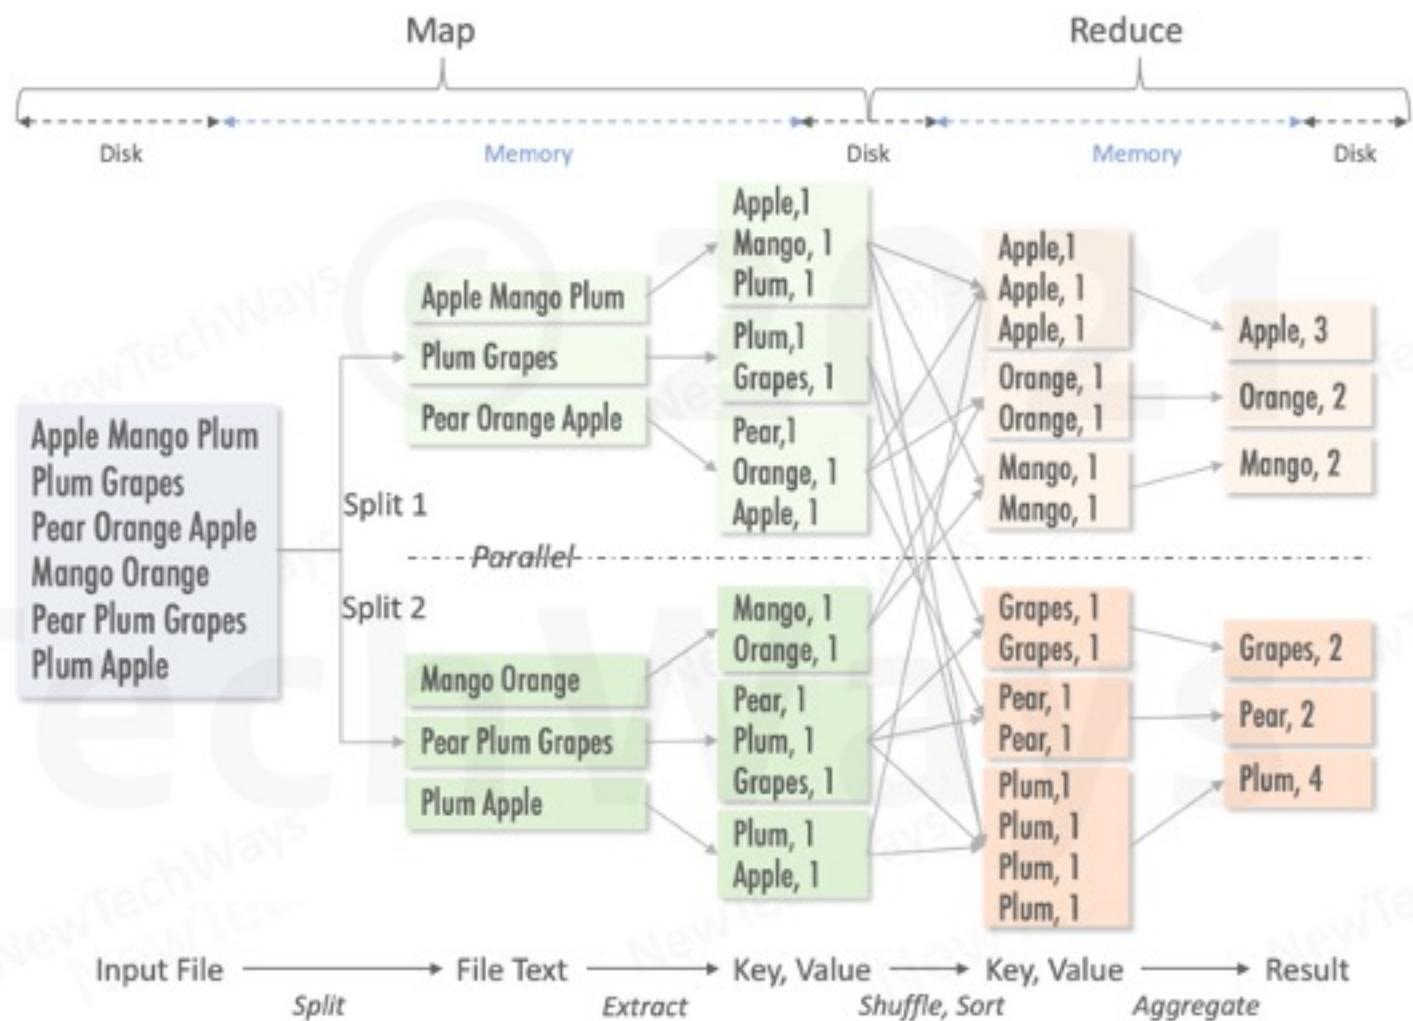
# Hadoop HDFS

- **Distributed File Data Storage**
  - Unstructured data files
  - Petabytes of data
  - Large file size > 100MB
- **Distributed files**
  - Files broken into chunks
  - For parallel reads
    - Map-Reduce
- **Sequential Writes – Append**
  - Large blocks of data – 64MB
- **Replication for reliability**



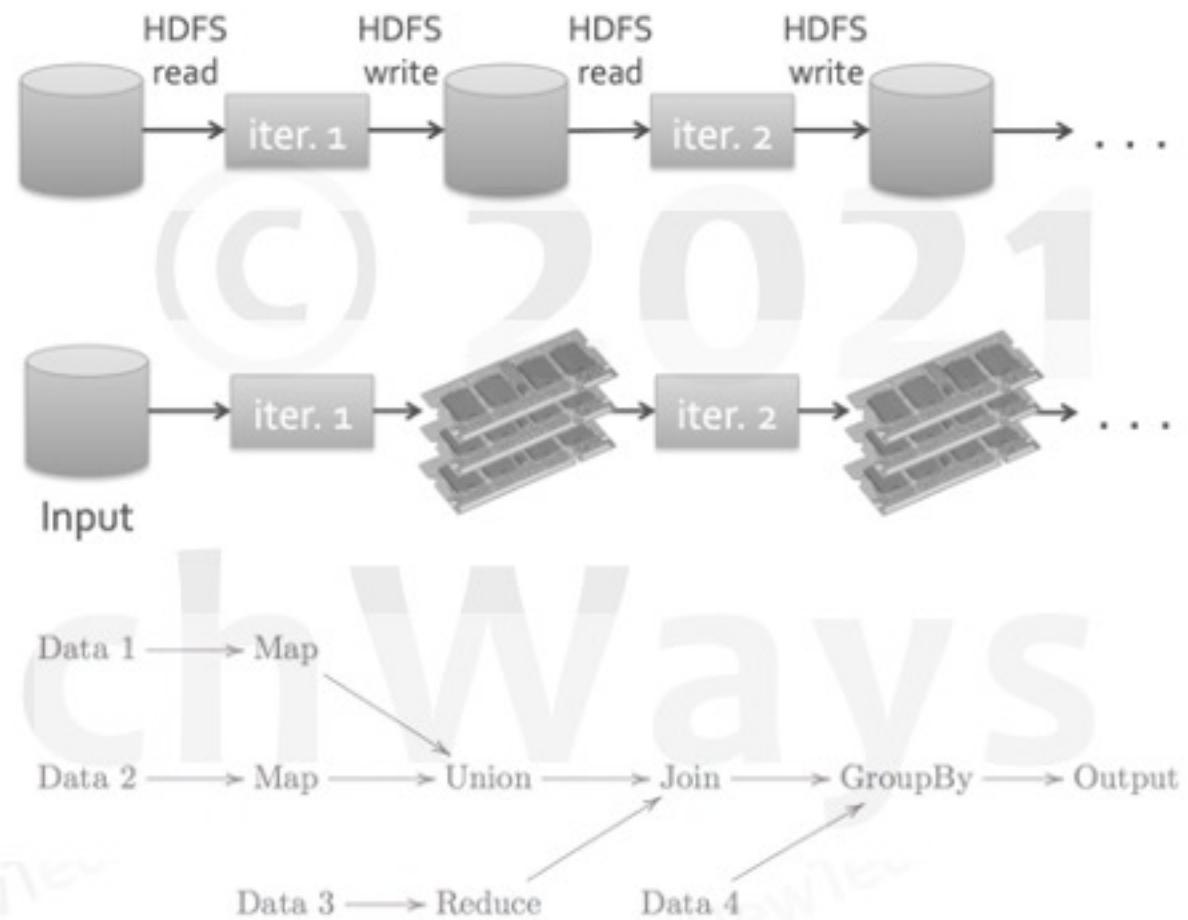
# Map-Reduce

- Parallel file processing on Hadoop cluster
  - Processing code executes on datanodes
  - Input/Output sources
    - HDFS, HBase, Cassandra
  - Map phase
    - Filtering and transformation
  - Reduce phase
    - Shuffles map output across nodes
    - Groups related information
    - Computes aggregate data



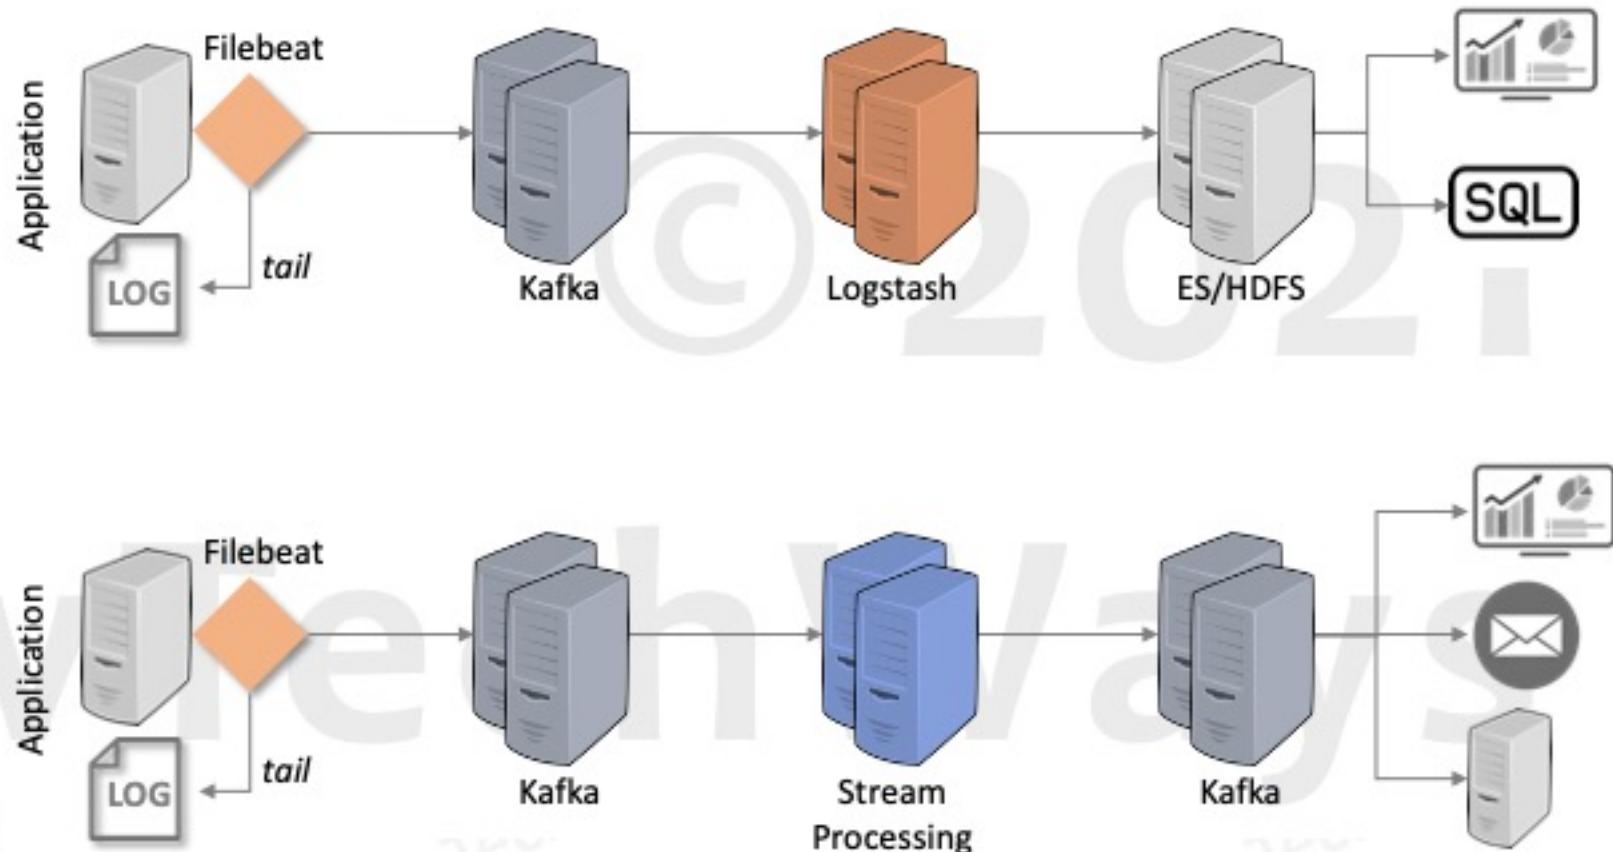
# Apache Spark

- Evolution of Map-Reduce
- In memory
  - 10x to 100x times faster
- DAG of operations
  - Multiple operations in a DAG
  - Multiple inbuilt operations
- Interactive
  - Interactive shell in Scala/Python/R
- In built libraries for
  - SQL Interface, Machine Learning, Graph Processing, Streaming



# Streaming Processing

- Low latency
  - Keep data moving
- High throughput
  - Multiple sources
- Event Issues
  - Event delay
  - Out-of-Order
  - Missing
- Fault Tolerance
  - Event Replay
- High Availability
- Processing Engine
  - Storm
  - Flink
  - Spark (Micro-Batching)
- Buffer
  - Kafka



Thanks!



**NewTechWays**

<https://www.newtechways.com>